

HPC

physical problem

↓
mathematic model

↓
numerical soln

↓
simulate

] HPC

performance metrics → power consumption, FLOPS

- ① Theoretical peak performance (R_{theory})
- ② Real performance (R_{real})
 - ↳ flops for specific operations.

sustained performance : for a particular period of time, total flops ($R_{\text{sustained}}$)

$$R_{\text{sustained}} < 10\% R_{\text{theory}}$$

$$R_{\text{sustained}} \ll R_{\text{real}} \ll R_{\text{theory}}$$

CPU not given data \Rightarrow CPU is idle

\rightarrow performance ↓

CPU \rightarrow ALU
 \downarrow CU

Dual core
 \rightarrow 2ALU
 \rightarrow 1CU

Quad core
 \rightarrow 4ALU
 \rightarrow 1CU

fast CPU requires fast memory

Computer time \rightarrow CPU cycle
 \downarrow min time to execute one instruction
Time to execute the programme

\rightarrow Time for 1 cycle \times no of instructions

Computer time measured in CPU cycle

Time to execute given programme min time to execute 1 instruction.

$$T = n_c \times t_c \Rightarrow n_i \times \frac{n_c}{n_i} \times t_c \Rightarrow n_i \times CPI \times t_c$$

$n_c \rightarrow$ no of CPU cycles
 $t_c \rightarrow$ time of the cycle

n_i = no of instructions

CPI \rightarrow no of clocks required for 1 instr

Reduce

1) cycle time (t)

increasing frequency

2) Reduce n_i (# of instr^{ns})

e.g.

$$c = a + b \rightarrow \text{inst 1}$$

\downarrow

load a

load b

Add a and b

store c



better

compile, Interpret

Assembler

③ reduce CPI

clock per instructions

parallelism & pipelining

Reduce CPI \rightarrow interleaved execution, overlapping execution, overlapping execution of # of instrn.

parallelism \rightarrow in single processor computer

① pipelining

② \uparrow no of cores

③ functional units

\rightarrow overlap operation of diff units of computer.

→ Increase speed of ALU by exploiting data temporal parallelism.

How to make computers fast.

Reduce CPI : Key is llism.

multiple CPU : → 2CPUs → llism in true sense
parallelism.

Simultaneous execution → llism.

Interleaved execution → ~~one task at a time~~ but CPU is never idle.

FLYN classification of 11 ARCH's.

If you are able to read / write 2 data items @ time → called llism.

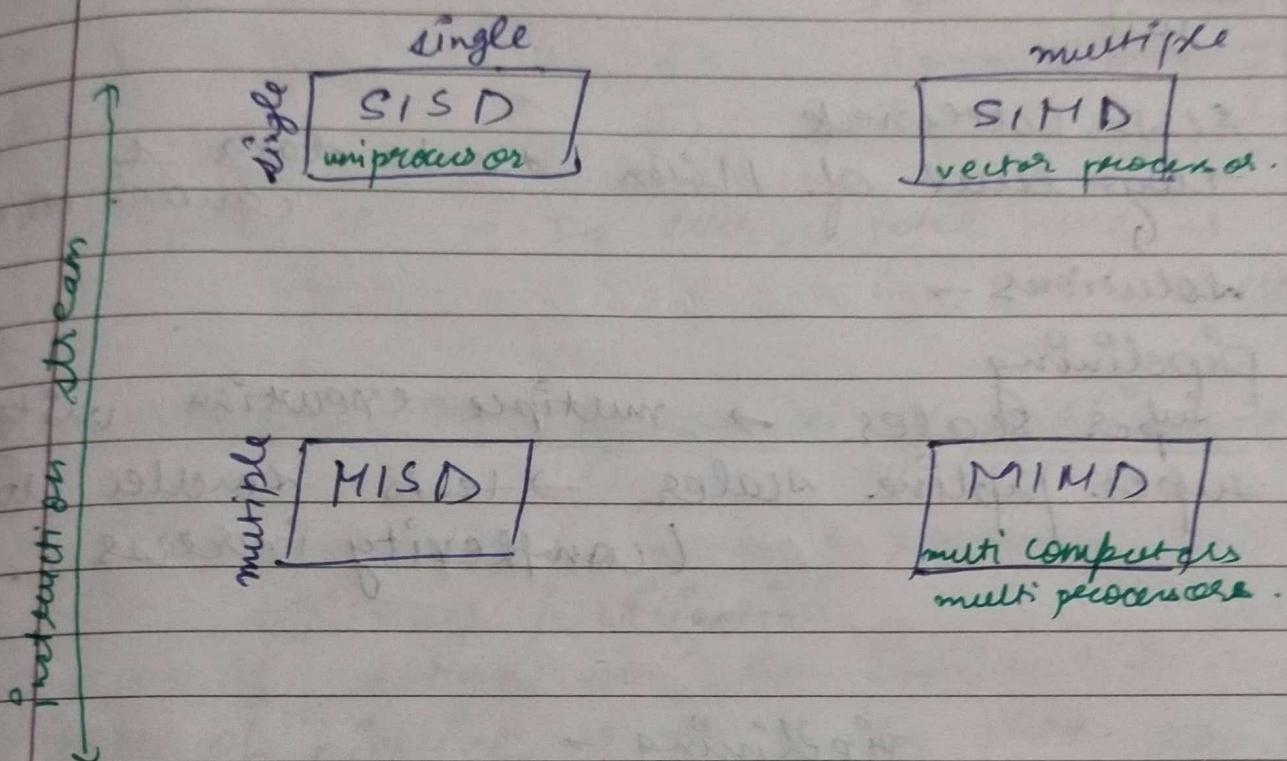
→ read / generation : etc.

collection of instructions → instruction stream.

data stream → { CPU → memory : write
memory → CPU : read

parallelism can happen on both data & instruction stream.

← data stream. →



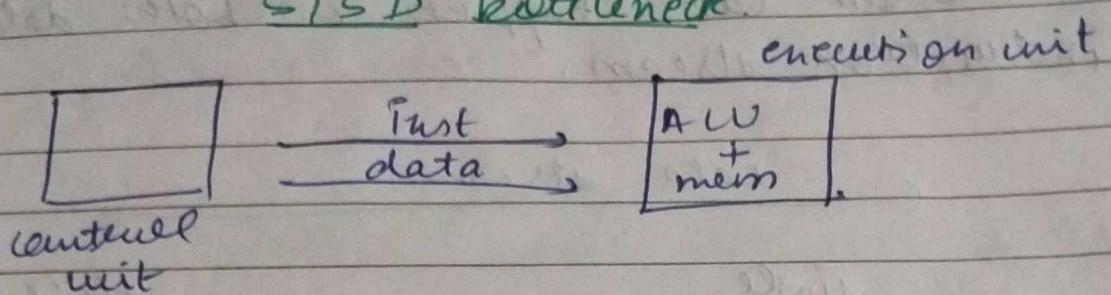
SISD : → Scalar processors → pipelining only solution to achieve 11xm here

→ Deterministic execution → tells what instructions are being executed when

SIMD → single instruction multiple data,

MISD → less significant.

MIMD → achieve 11xm in true sense.

SISD bottleneckSISD bottleneck

→ low level of I/Oism : due to data & control dependency

solutions →

- ① Pipelining
- ② Super scalar → multiple execution units
- ③ super pipeline scalar → fetches smaller chunks
(complexity increase).

Pipelining →

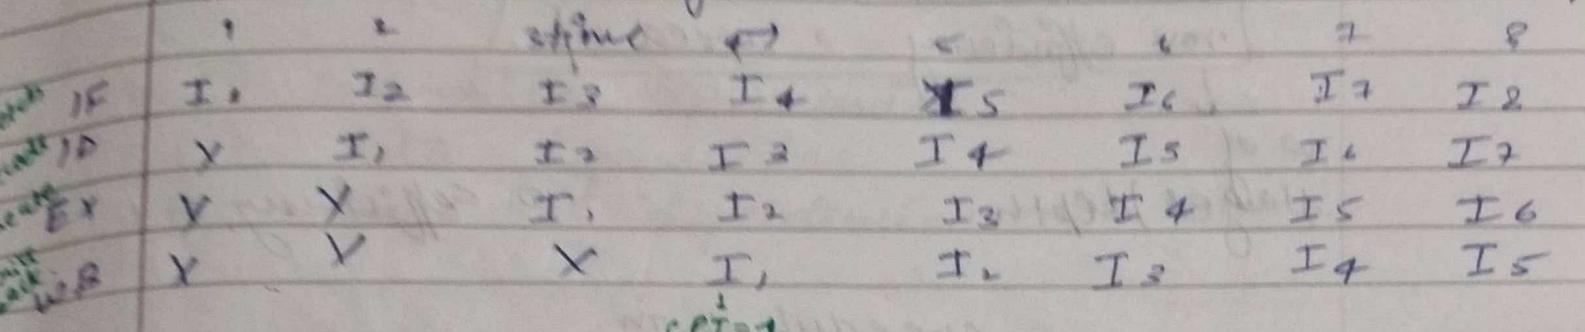
2 instruction per cycle → goal

→ overlapped / interleaved execution → pipelining.

→ 1 instruction can be divided into 4 phases of stages

- i) Instruction fetch → memory to register
- ii) Read/decodes → memory → register.
- iii) execution → ALU
- iv) write back → reg → memory

for 7 instructions → $3 \times 4 = 12$ cycles.
 without pipelining → $3 \times 7 = 21$ cycles (7 instructions)



after 7 → I7 over total time → 7.
 ↳ after pipelining

measure the performance of pipelines →

- ① speedup
 - ② efficiency justification
- no of stages in an instruction.
- depth of pipeline → no of ops that can be performed in 1 cycle.

$$CPI = \frac{n + \text{depth} - 1}{n} \Rightarrow \frac{\text{total no of CPU cycles}}{\text{total no of instructions.}}$$

efficiency → $\frac{\text{no of boxes utilised}}{\text{Total no of pipelines.}}$

speed up → $\frac{\text{Time to execute without pipelining}}{\text{Time with pipelining}}$

→ ~~Eff.~~ ~~Eff.~~ ~~Eff.~~

→ in general -

$n \times \text{efficiency} \rightarrow \text{speed up.}$

$\frac{\text{no of stages}}{\text{depth}}$

○ efficiency < 1

speedup < m.

$$\text{speedup} = \frac{D}{CPI} \rightarrow \text{depth.}$$

$$T_{\text{nop}} \rightarrow \text{Time} \Rightarrow 1 \times \text{depth}/m + (n-1) \times 1$$

without pipeline

m → depth.

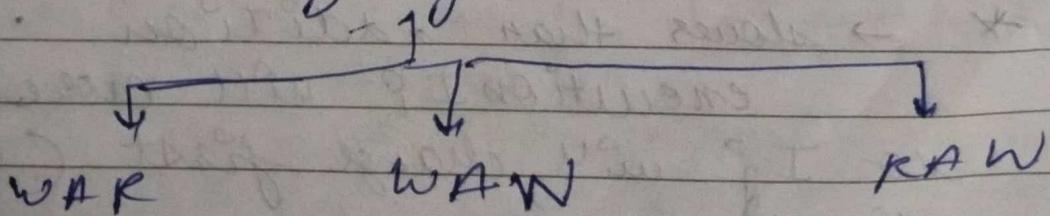
, n → no of instructions

$$T_{\text{wp}} \quad \text{Time with pipeline} \Rightarrow nxm.$$

$$\text{now speedup} \rightarrow \frac{T_{\text{nop}}}{T_{\text{wp}}} \Rightarrow \frac{nxm}{m+n-1}$$

Inhibitors of pipelining

- ① control dependencies
 - ② data dependencies
 - ③ structural dependencies
- due to branch statements. → branch penalties happen.
 when next operation is jump statements.
 Then jump should run first we have to wait for the upper instruction & not start new
 ↳ use NOP operation
 ↳ done by smart compilers
 ↳ not hardware support
- pipelining + stalls → Hardware support
- ↓
 different from NOP.
- rearrange instructions → software
 → rescheduling instructions
 → Branch prediction → hardware
 → Data dependency →
- when an instruction depends on the data coming from other instructions.



RAW → read after write [flow] true data dependency

WAW → write after read [anti-data dependency]

WAW → write after write [output data dependency]

examples

RAW: read after write

$$I_i: R_1 \leftarrow R_1 + R_3$$

$$I_j: R_4 \leftarrow R_1 + R_5 .$$

WAR:

$$I_i: R_1 \leftarrow R_2 + R_3 \quad R_2 \text{ value vs to}$$

I_j : load m, R_2 read here
its value

I_j is update earlier

WAW: write after write [output data dependency]

$$I_i: R_1 \leftarrow R_2 * R_3$$

$$I_j: R_1 \leftarrow R_5 + R_6$$

* → slower than addition

execution I_j will more than I_i
so I_j will change first (execute first).

sequence not followed.

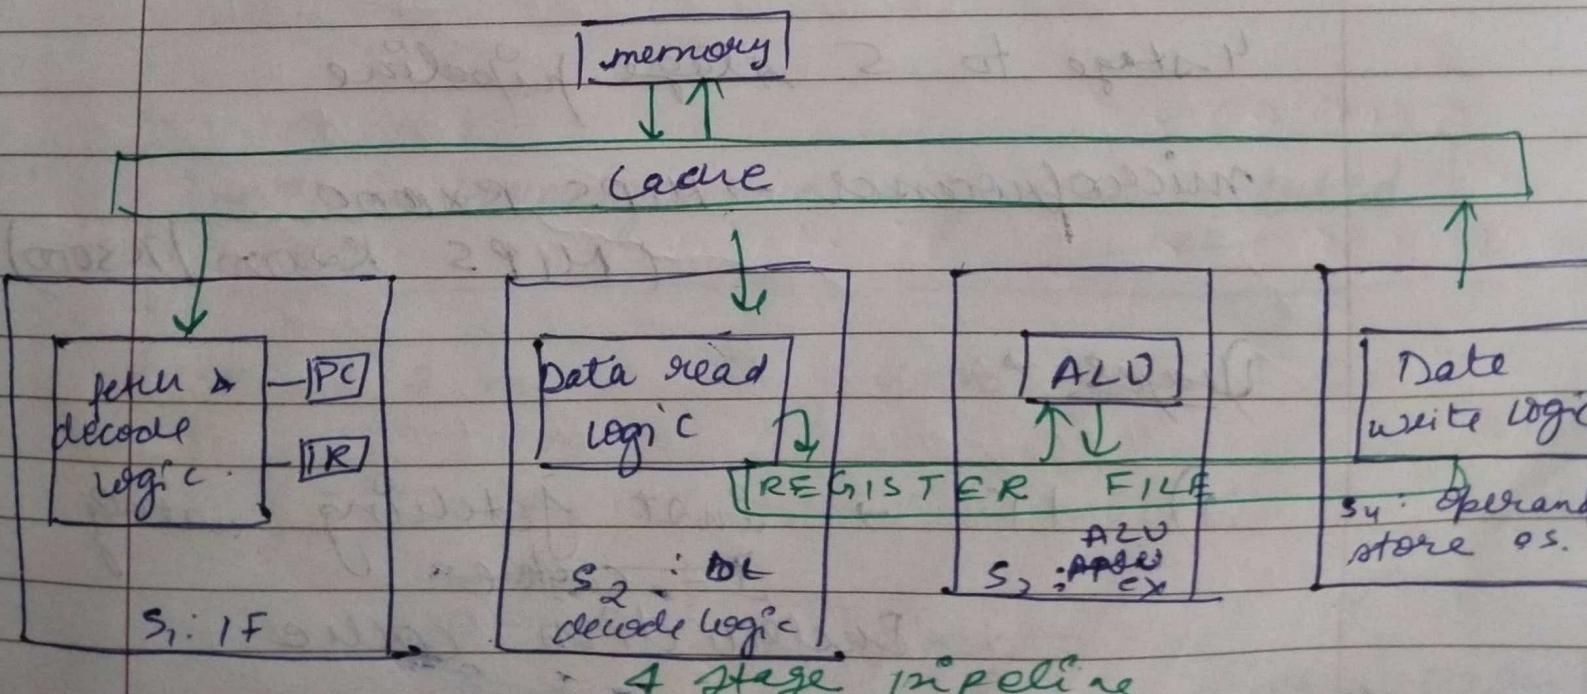
solutions →

lock → draback → stall
stall → no op.

Integrate registers → hold intermediate outputs of p.
↳ called operand forwarding

structural dependence →

resource conflict in the pipeline
eg → Bus
when more than one instruction tries to fetch the same resource in same cycle then structural dependence.



same time S₁ & S₂ can ask for cache conflict

I1	1	12	3	12
I2	2	2	2	3
I3	1	2	1	2

Page No.	
Date	

sol" →

- ① resource duplication (instead we have multiple copies of resource)

here divide cache into 2 parts

① instruction cache

② data cache

problem 2

instruction 3 at S_2 & 1 at C_2 both accessing
at time T_4

solution →

4 stage to 5 stage pipeline

microprocessor → MIPS RX1000
(MIPS R2000/R3000)

stages →

- 1) IF → first fetching using
~~I-cache~~
Instruction cache

- 2) RD : operand loading (reading) from register file 'RF' will be decoding the fetched instruction.. whatever operands you require are loaded into register file .
- 3) execute → data processing using ALU & register file . as needed
- 4) memory access → operand accessing (load / store) using data - cache
- 5) write back → operand storing (written back in register file).

types of bus →
now,

$$CPI = 1$$

what is throughput → no of inst executed
2 sec

$CPI = 1 \rightarrow$ means in 1 clock cycle, 1 inst is execute

if CC = 1 ns \rightarrow 1 ns , 1 unit is execute

$$10^{-6} \text{ sec} = 1 \text{ inst}$$

$$1 \text{ sec} \rightarrow \frac{1}{10^{-6}} \text{ inst executions}$$

Types of pipeline

① Uniform delay pipeline

② non-uniform delay pipeline

every stage will take same time to complete an operation

normal cycle time (T_p) = stage delay

cycle time (T_p) \rightarrow stage delay + buffer delay
after stall adding

if we have \rightarrow stages let's say

use 3sec 2sec.

all will be given 4 sec \rightarrow now

&

add stalls.

non-uniform delay \rightarrow diff stage

Cycle Time (T_p) = max(stage delay)

for
uniformity

buffer

$T_p = \max(\text{stage} + \text{Buffer delay})$

not necessary that after buffer all are same.

SIMD

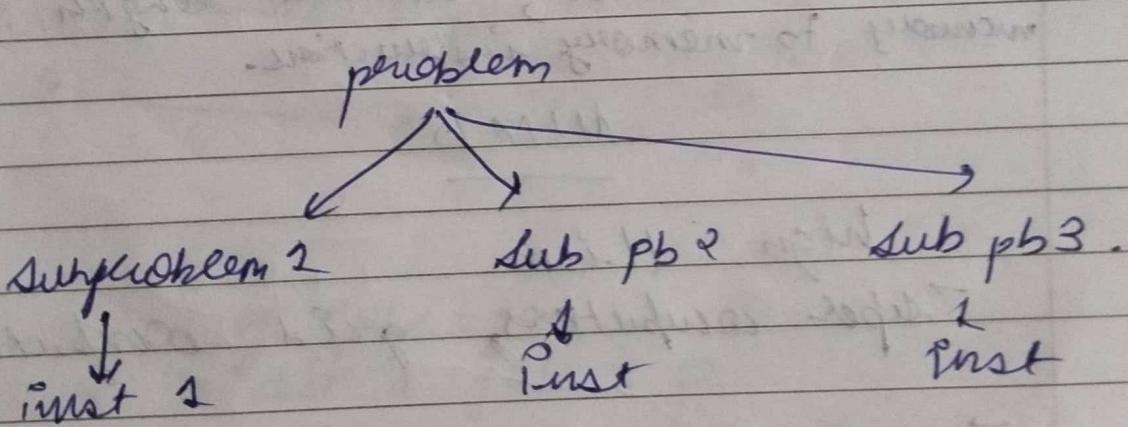
instruction

same → data different

deterministic data → we know the output

synchronous $s_1 \rightarrow s_2 \rightarrow s_3$
behavior ↴

↳ data source

in SIMD should know at what time
which data is getting from where
eg at $T_1 \rightarrow$ data from s_1 .

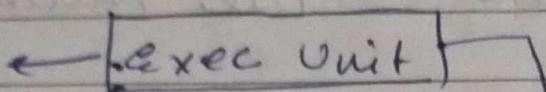
eg → addition is a subproblem and
multiple data → $(1, 2), (3, 4)$ are
coming
so instruction → single
data → multiple.

MISD

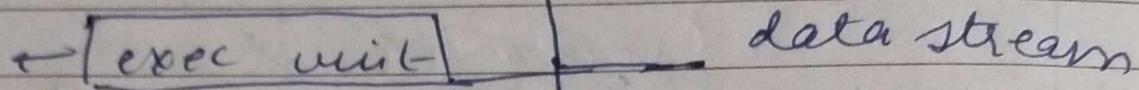
MISD bottleneck

- low level of 11.1nm.
- high sync required synchronization.

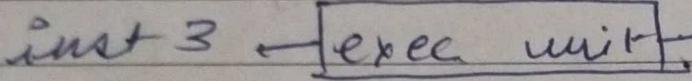
inst 1



inst 2



inst 3



bottleneck.

- bottleneck in memory access, microcode, resource contention, instruction decoding.

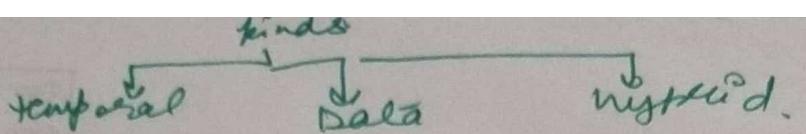
- CISC → complex instruction set computers
- RISC Instruction set, variable length instruction, memory to memory instructions.

MIMD

high 11nm

super computers, grid computers

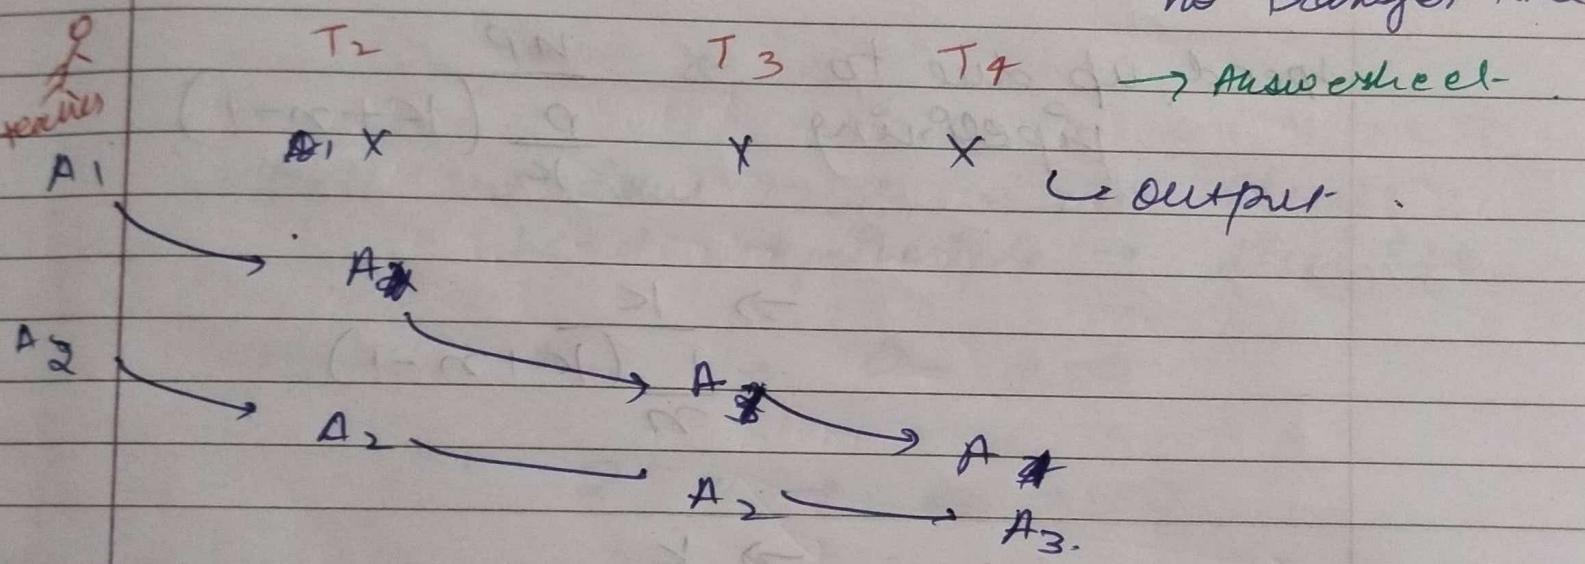
- no sync necessarily required but can be both.
- no data is shared → loosely coupled.
- if sharing then sync.



GPU → graphical processing units.
specific type of test is assigned to small components.

Kinds of parallel alarm

→ temporal / test → all tests should be similar with no delay
→ no transfer time



let the no of jobs = n
let time to do a job = p

let each job be divided into k tasks &
let each task be done by a different individual

let the time for each task = $\frac{p}{k}$.

time to complete → np .
n jobs with no pipelining.

with a pipeline organisation of K individuals

$$\Rightarrow P + (n-1) \frac{P}{K}$$

$$\Rightarrow \frac{KP + (n-1)P}{K}$$

$$\Rightarrow P \frac{(K+n-1)}{K}$$

speed up due to pipelining $\rightarrow \frac{NP}{\frac{P}{K}(K+n-1)}$

$$\Rightarrow \frac{K}{\frac{1}{n}(K+n-1)}$$

$$\Rightarrow \frac{K}{(1 + \frac{K-1}{n})}$$

problems to temporal alignment

\rightarrow Transfer time \rightarrow synchronisation

we assumed this negligible.

- ② fault tolerance (all processes are equal)
one dies fault can't be solved/fixed.
- ③ scalability not possible
- ④ inter task comp communication
- ⑤ bubbles in pipeline (one stops all other stops)

Data 11 am.

suppose we have 1800 answer sheets.
read → principle → divides.

teacher one → 0 - 250

" 2 → 250 - 500

" 3 → 500 - 750

" 4 → 750 - 1000

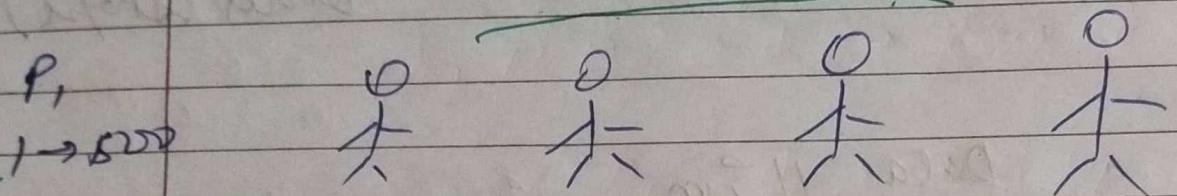
→ so all teachers are assigned tasks at one time together & again assigned together → (no overhead, inter process communication)

disadvantage →

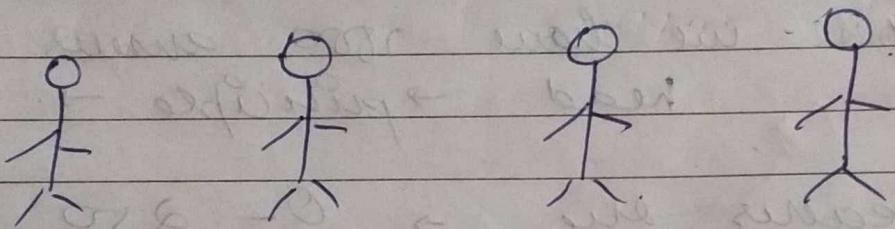
→ single point of failure.

→ task should be ~~mutually~~
 → cost more (need takes more salary
 before any 4 teachers).

Mixed Job



[data Job]



eg \rightarrow NTC-SX
 OFF. CRAY

Data Job \rightarrow with dynamic assignment

instead giving 1-250 at same time
 give next if previous job is done by the process or. (1 given at a time)

Quasi - dynamic \rightarrow now one time in
 dynamic giving 4-5 at a time.

specialist data II ism →

specialist person → who knows who can do what?

1. cause Grained specialist temporal parallel processing

in ways
out ways of the processor.

Detecting II ism

→ can this run on a II arch

if p. Pism.

dependent → cannot

independent → can run.

Bernstein's condition we define →

$I_i^r \rightarrow$ input set → from have to fetch
read from memory

$O_i^w \rightarrow$ represent write set → write into
memory

process →

read set

write set

P_1 and P_2 will be independent when

$I_1 \cap D_2 = \emptyset$ only when they are independent.

$I_1 \cap D_2 = \emptyset \rightarrow$ flow independence

$I_2 \cap D_1 = \emptyset \rightarrow$ anti independence

$D_1 \cap D_2 = \emptyset \rightarrow$ output independence

eg →

$$P_1: C = D \times E$$

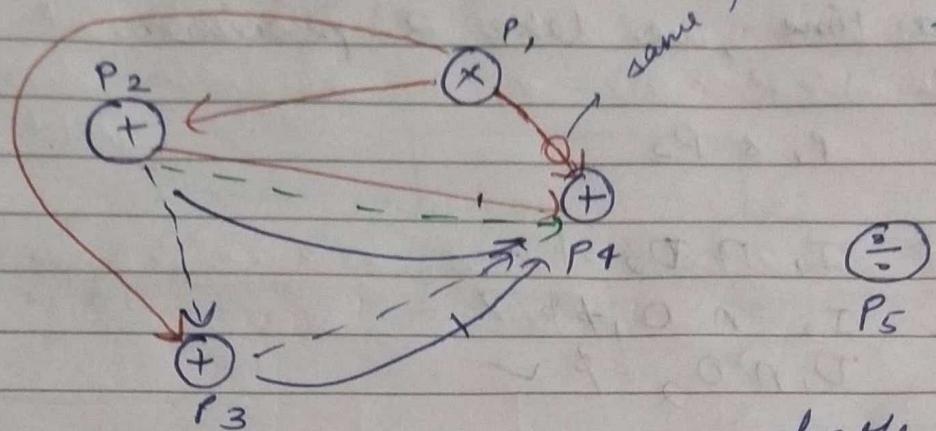
$$P_2: M = G + C$$

$$P_3: A \Rightarrow B + C$$

$$P_4: C = L + M$$

$$P_5: F = G \circ E$$

dependence graph \rightarrow

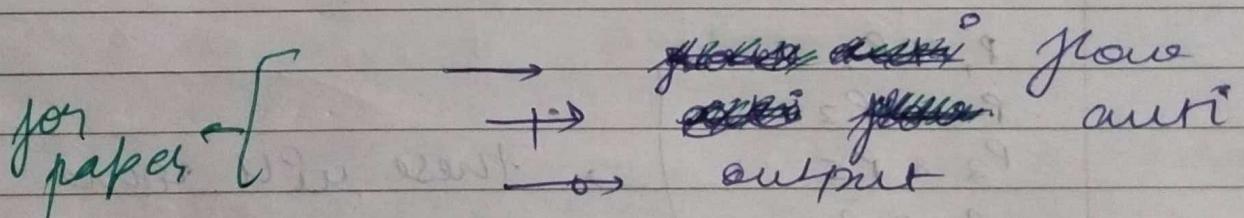


both use +
so resource
dependency

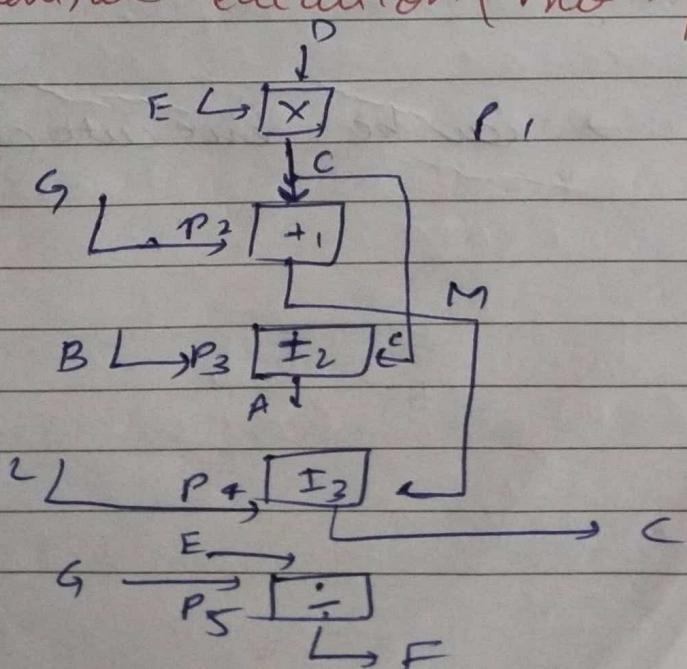
there are 3 addition units ($+, +_2, +_3$).

Solid arrows \rightarrow data dependence.

Dotted arrows \rightarrow resource dependence.



sequential execution (no pipelining)



II execution →

at a time, we take 2 processes

$P_1 \& P_2$

$I_1 \cap D_2 = \emptyset \checkmark$

$I_2 \cap O_1 = \emptyset \times$

$D_1 \cap O_2 = \emptyset \checkmark$

for pairs (all) when all 3 conditions are true then they can run parallelly.

$P_1 \& P_5$

$P_2 \& P_3$

$P_2 \& P_5$

$P_3 \& P_5$

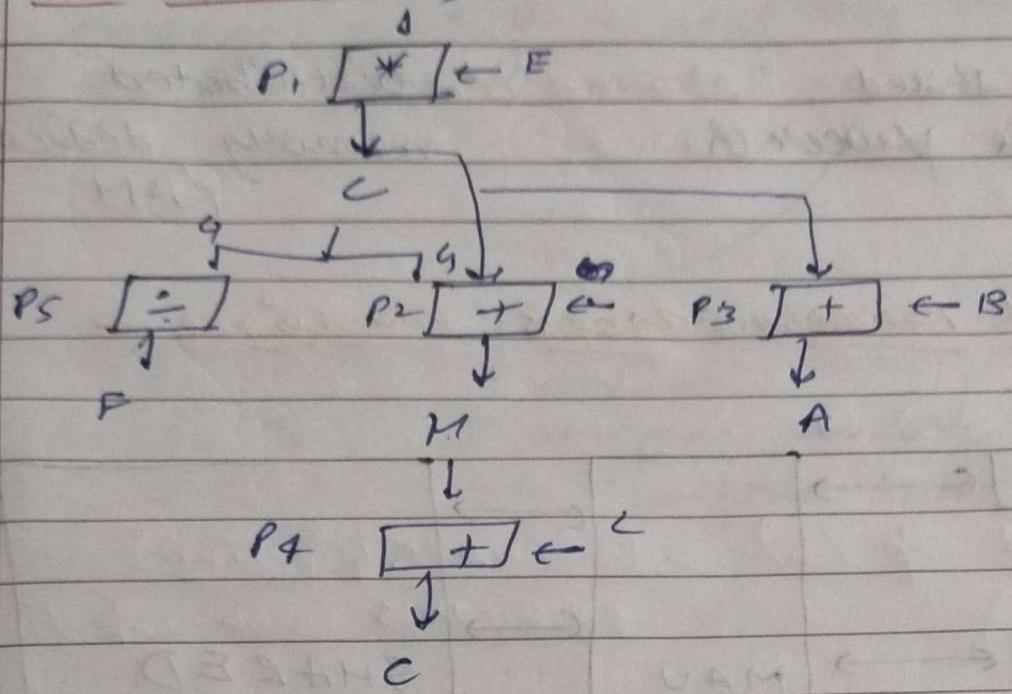
$P_4 \& P_5$

these will come II.

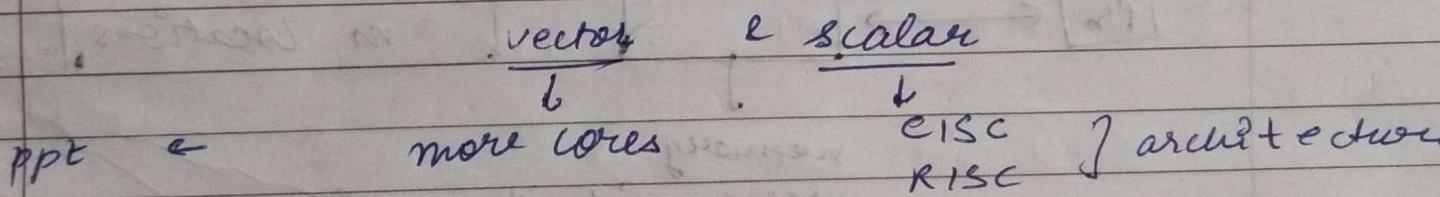
$P_2, P_3, P_5 \rightarrow$ can be executed II.

Q

11 execution



Processor types



locality of reference → Basic principle on which cache works.

→ spatial / temporal locality

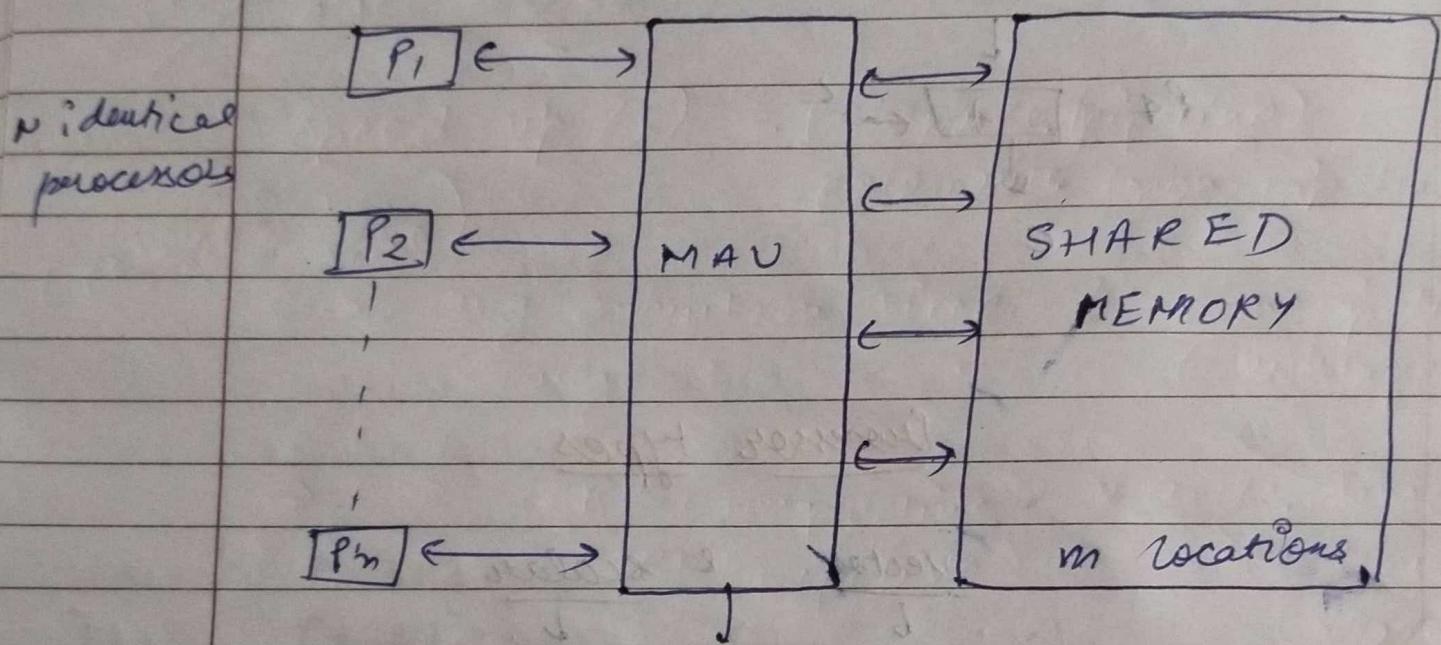
→ cache

→ pipelining

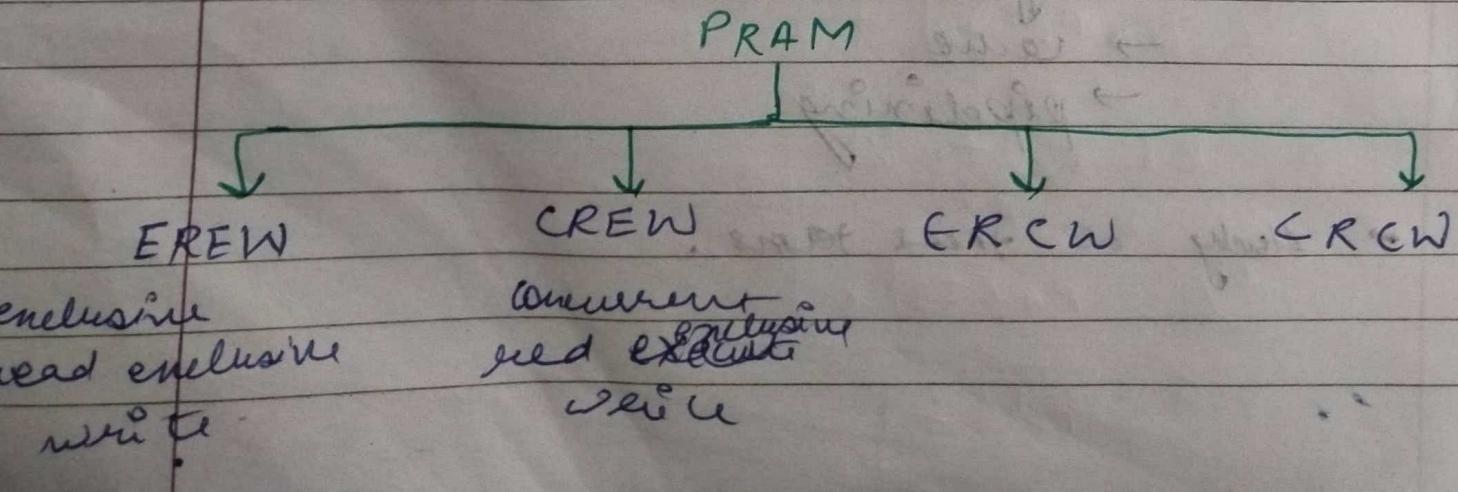
→ Study all these terms.

multicore processors and multicore computers

(PRAM) parallel random access machine →



all n processors has equal access to the shared memory.



- ① shared memory model.
- ② message passing model.
- ③ data parallel model.

shared memory

uniifrom

- ① UMA \rightarrow uniform memory access
- ② NUMA \rightarrow non-uniform memory access
- ③ cache only

$P_1 \rightarrow [SM1]$

$P_2 \rightarrow [SM2]$

$P_3 \rightarrow [SM3]$

\rightarrow each processor has its own memory and but other ~~processor~~ can access the other's memory equal.

NUMA \rightarrow

every P will have local RAM

$[P_1] \rightarrow LR$ (local Ram)

$[P_2] \rightarrow LR$

$[P_3] \rightarrow LR$

eg \rightarrow super computer

here P_1 has more right on LR
 (all do not have same right)
 processors are loosely coupled.

Cache only (C0NA) →

only cache shared b/w processors.
RAM not shared.

~~parallel programming paradigms~~ →

slide 3 →

$$S \rightarrow \frac{\text{speed up (s)}}{\text{performance (non-pipeline)}} = \frac{\text{Performance (pipeline)}}{\text{Performance (non-pipeline)}}$$

(Response & $\frac{1}{\text{Time}}$)

$$S \rightarrow \frac{\text{avg execution time (non-pipeline)}}{\text{avg. " " (with pipeline)}}$$

• $S \Rightarrow \frac{\text{CPI (non-pipeline)}}{\text{CPI (pipeline)}} \times \frac{\text{cycle time (non-pipeline)}}{\text{cycle time (with pipeline)}}$

due to stalls →

$$S \Rightarrow \frac{\text{CPI non pipeline}}{(1 + \text{no of stalls per instruction})}$$

because ideally 11.1cm CPI ⇒ 1