HTTP & REST

COCSC20

What is Cloud Computing?

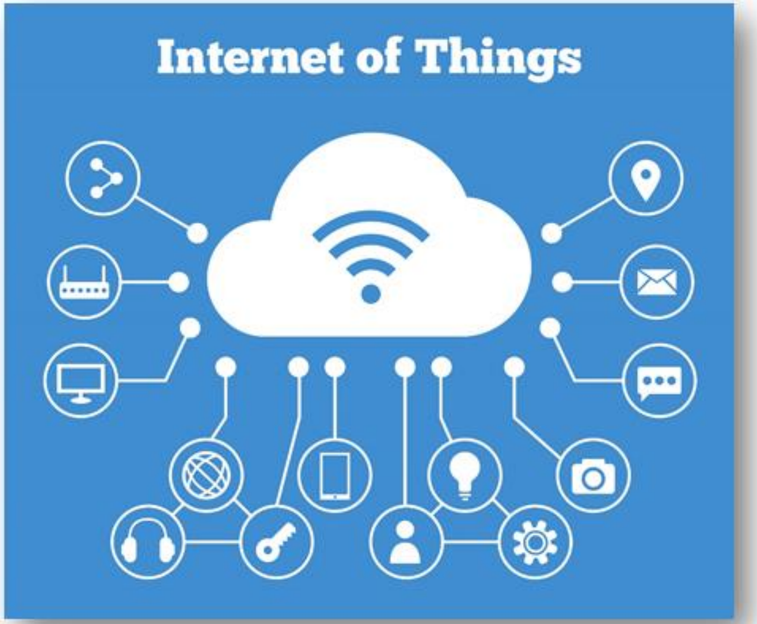

What is Cloud Computing?

- Delivery of computing **services**
  - servers
  - storage
  - analytics
  - databases
  - networking
  - and much more...

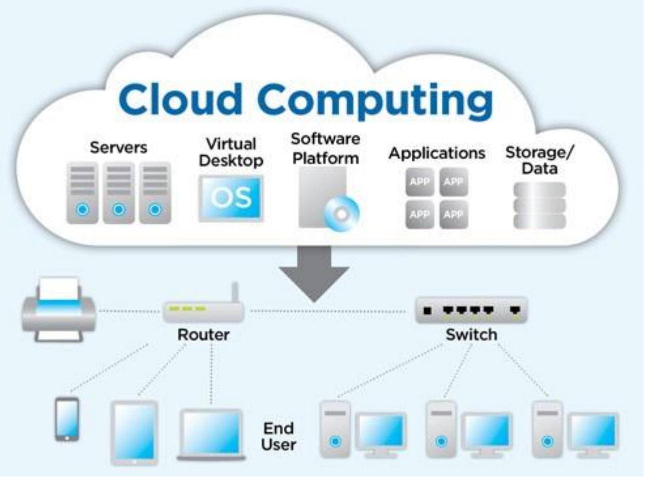

- Another definition: network-based computing taking place over the Internet, while hiding complexity of underlying infrastructure using simple APIs.

What is Cloud Computing?

- Collection/group of **integrated and networked hardware, software, and Internet infrastructure** (called a platform)

- Platforms provide **on demand** services that are **always on** and **accessible anytime** and **anywhere**

What is Cloud Computing?

- Advantages:
  - New applications
  - Anytime/anywhere access
  - Homogeneity
  - Virtualization
  - Resilient
  - Cost
  - Sharing, collaboration
  - Management/maintenance
  - Security
  - …

Cloud Models: IaaS, PaaS, SaaS

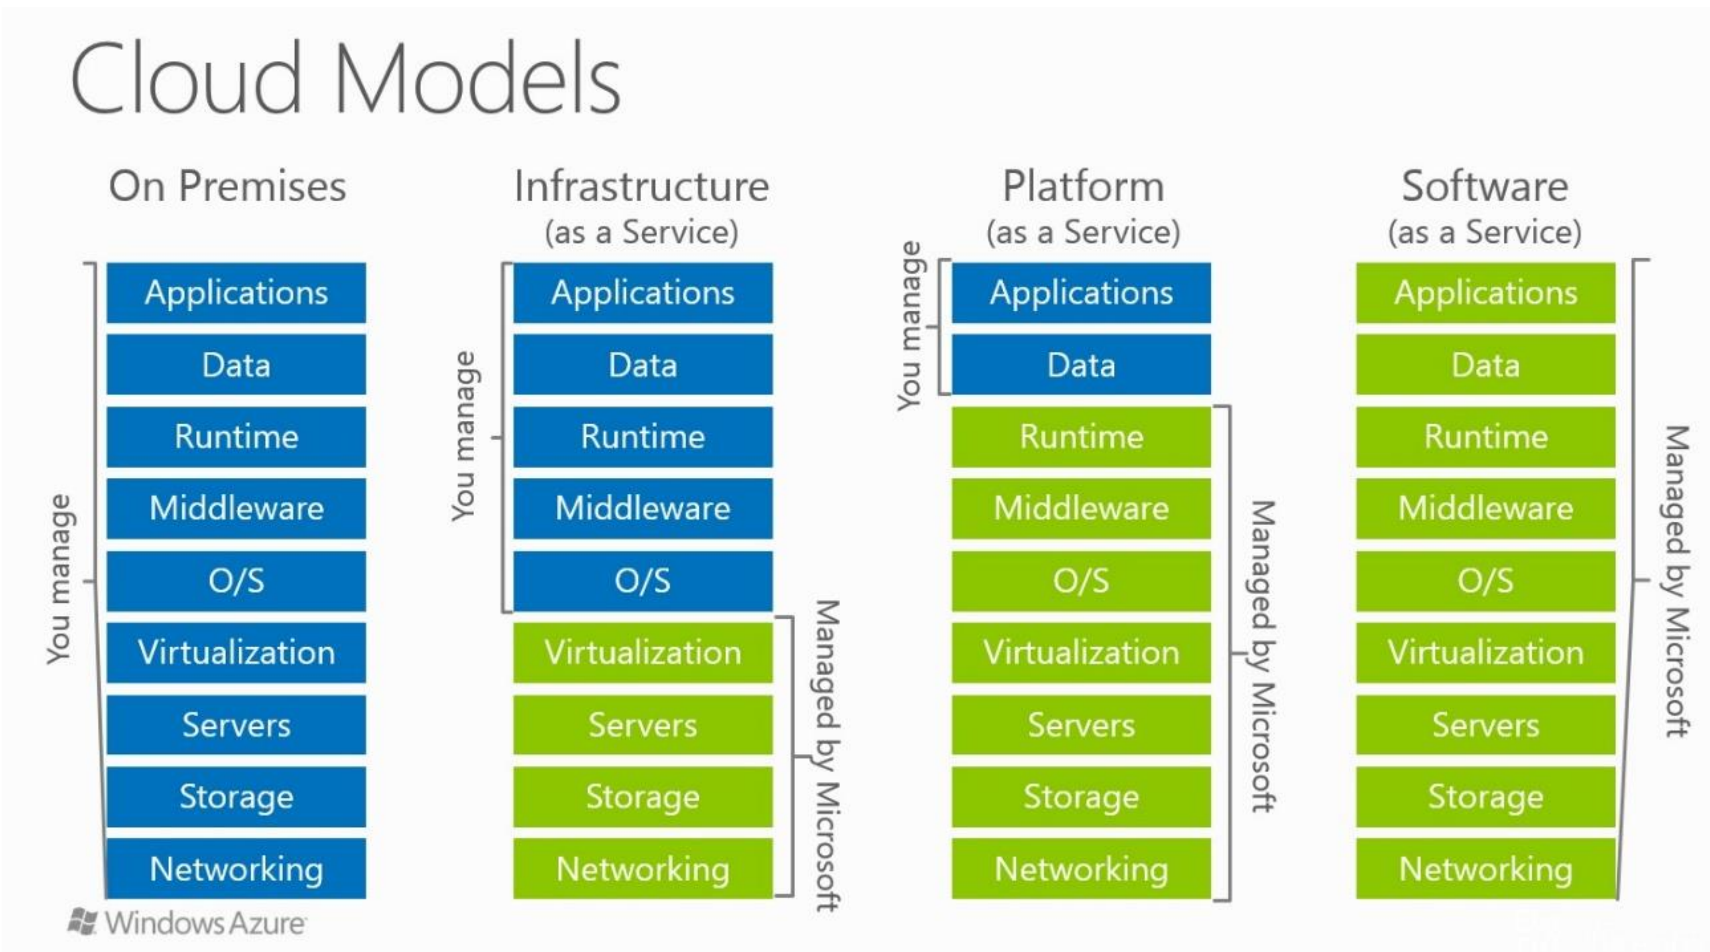

Definitions

- **Virtualization:** creation of a virtual resource such as a server, desktop, operating system, file, storage, or network
- **Middleware:** software that acts as a bridge between an operating system or database and applications, especially on a network
- **Runtime:** software designed to support the execution of computer programs

IaaS, PaaS, SaaS

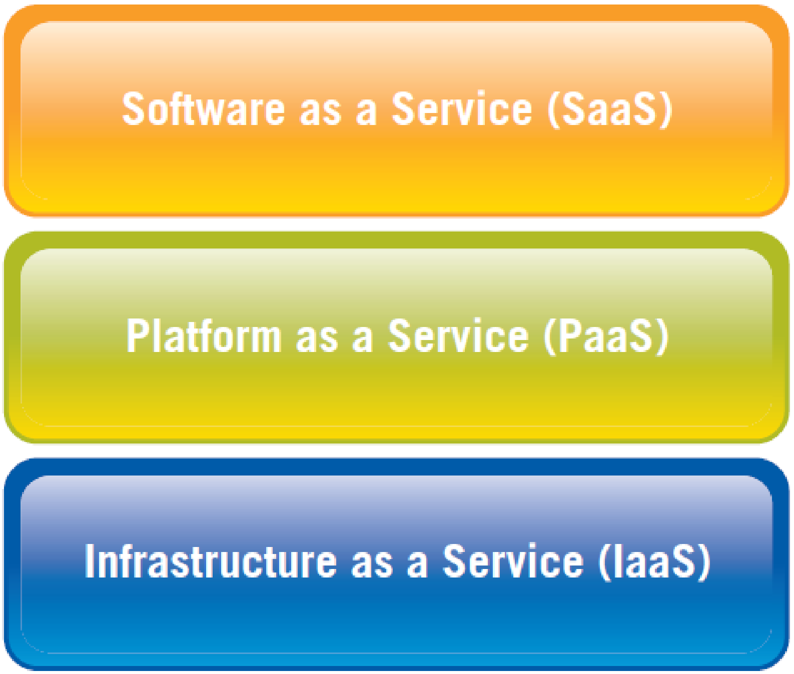

Enduser application is delivered as a service. Platform and infrastructure is abstracted, and can deployed and managed with less effort.

Application platform onto which custom applications and services can be deployed. Can be built and deployed more inexpensively, although services need to be supported and managed.

Physical infrastructure is abstracted to provide computing, storage, and networking as a service, avoiding the expense and need for dedicated systems.

Simple example:

- IaaS: barebones computer
- PaaS: computer + OS (incl. development environment)
- SaaS: complete solution including application(s)

IaaS, PaaS, SaaS

- IaaS: Amazon Web Services (AWS), Microsoft Azure, Google Compute Engine
- PaaS: Google App Engine, Heroku, OpenShift, AWS Elastic Beanstalk
- SaaS: Google Apps, Dropbox, Cisco Webex, Salesforce, Concur, GoToMeeting

Basic Cloud Characteristics

- “No-need-to-know”: interact with underlying infrastructure via API
- Flexibility and elasticity: scale systems up and down (allocate/release resources) based on needs
- Pay as much as used and needed (actual usage vs. service levels)
- Anytime anywhere access

Virtualization

- Virtual workspaces:
  - An abstraction of an execution environment that can be made dynamically available to authorized clients by using well-defined protocols
  - Resource quota (e.g., CPU, memory share)
  - Software configuration (e.g., OS, provided services)
- Implemented on Virtual Machines (VMs):
  - Abstraction of a physical host machine
  - Hypervisor intercepts and emulates instructions from VMs, and allows management of VMs
  - VMWare, Xen, etc.

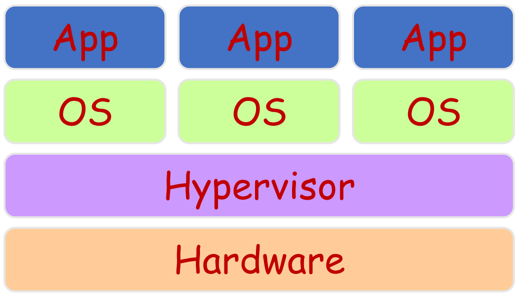

Virtualized Stack

Virtual Machines

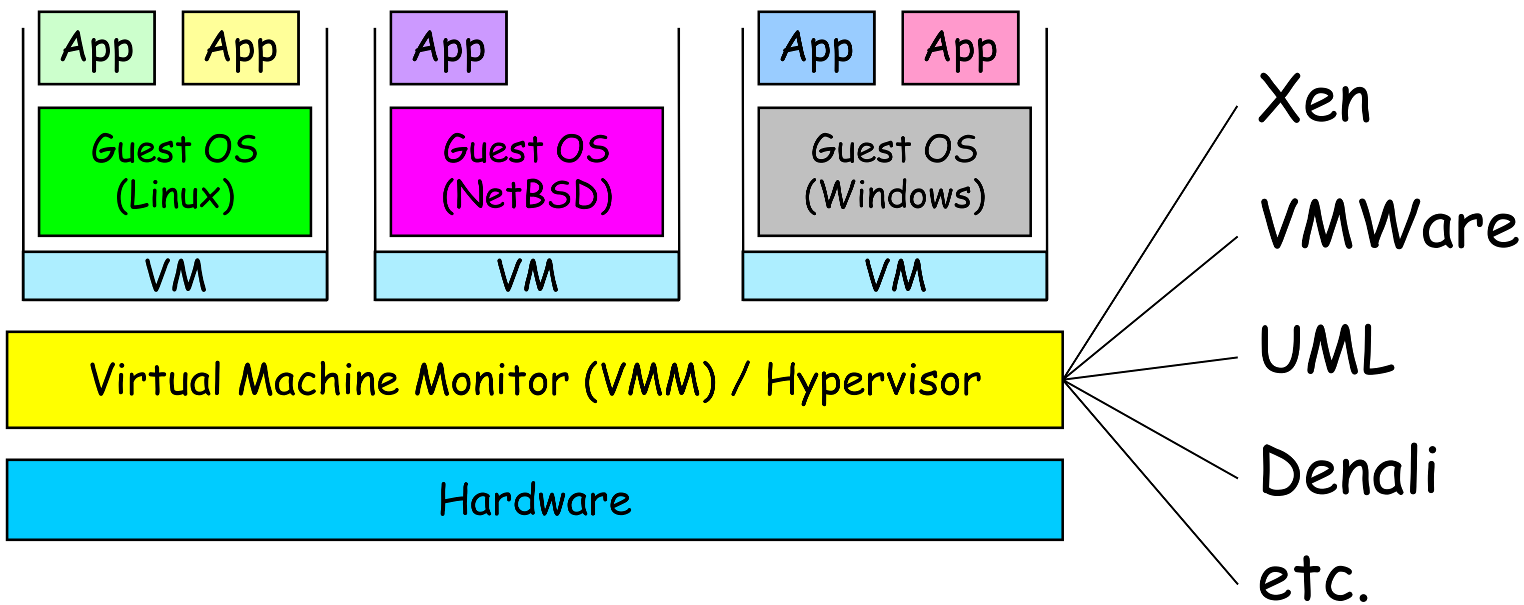

Cloud Example: S3

- Amazon Simple Storage Service (S3)
- Unlimited storage
- Pay for what you use

	S3 Standard	S3 Standard – Infrequent Access	AWS Glacier
STORAGE			
First 50 TB/ month	$0.023 / GB	$0.0125 / GB	$0.004 / GB
Next 450 TB/ month	$0.022 / GB	$0.0125 / GB	$0.004 / GB
Over 500 TB/ month	$0.021 / GB	$0.0125 / GB	$0.004 / GB
REQUESTS			
PUT, COPY, POST, or LIST	$0.005 / 1,000 requests	$0.01 / 1,000 requests	
GET and all other requests	$0.004 / 10,000 requests	$0.01 / 10,000 requests	
Delete requests	Free	Free	Free, but with limits and potential surcharges
Lifecycle Transition Requests into S3 Standard IA		$0.01 / 1,000 requests	
Glacier archive and restore requests			$0.05 / 1,000 requests, see Glacier pricing for more details on retrieval fees

Cloud Example: EC2

- Amazon Elastic Compute Cloud (EC2)
  - Virtual computing environments (“instances”)
  - Pre-configured templates for instances
  - Launch as many virtual servers as needed (“elastic”)
  - Xen and KVM hypervisor

Do You Use The Cloud?

Cloud for IoT

AWS IoT

ThingSpeak

Samsung ARTIK™ Cloud

Microsoft Azure IoT Platform

IBM Watson IoT™

HyperText Transfer Protocol (HTTP)

- Clients and servers communicate using the **HyperText Transfer Protocol (HTTP)**
  - Client and server establish TCP connection
  - Client requests content
  - Server responds with requested content
  - Client and server close connection (usually)

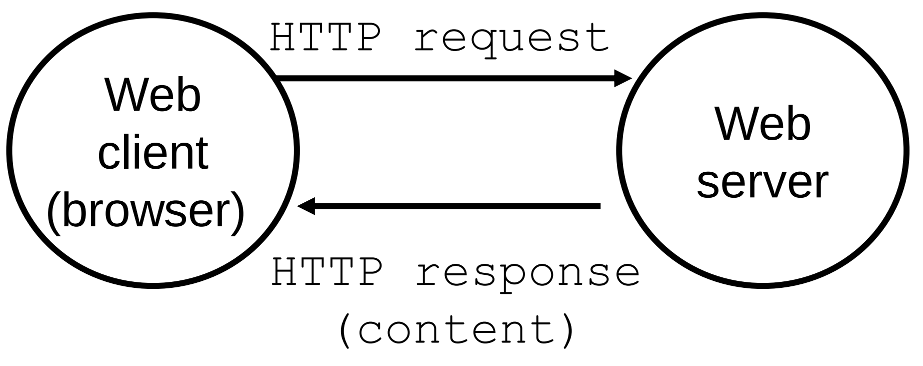

Web Content

- Web servers return **content** to clients
  - a sequence of bytes with an associated MIME (Multipurpose Internet Mail Extensions) type
- Example MIME types
  - `text/html` HTML document
  - `text/plain` Unformatted text
  - `application/postscript` Postcript document
  - `image/gif` Binary image encoded in GIF format
  - `image/jpeg` Binary image encoded in JPEG format

Static & Dynamic Content

- The content returned in HTTP responses can be either **static** or **dynamic**
  - Static content: content stored in files and retrieved in response to an HTTP request
    - Examples: HTML files, images, audio clips
  - Dynamic content: content produced on-the-fly in response to an HTTP request
    - Example: content produced by a program executed by the server on behalf of the client
- Bottom line: all web content is associated with a **file** that is managed by the server

URLs

- Each file managed by a server has a unique name called a URL (Universal Resource Locator)
- URLs for static content:
  - `http://www.cse.nd.edu:80/index.html`
  - `http://www.cse.nd.edu/index.html`
  - `http://www.cse.nd.edu`
    - Identifies a file called `index.html`, managed by a web server at `www.cse.nd.edu` that is listening on port 80
- URLs for dynamic content:
  - `http://www.cse.nd.edu:8000/cgi-bin/adder?15000&213`
    - Identifies an executable file called `adder`, managed by a web server at `www.cse.nd.edu` that is listening on port 8000, that should be called with two argument strings: `15000` and `213`

Anatomy of an HTTP Transaction

`<unix>` *`telnet www.aol.com 80`*	*Client: open connection to server*
`Trying 205.188.146.23...`	*Telnet prints 3 lines to the terminal*
`Connected to aol.com.`	
`Escape character is '^]'.`	
`GET / HTTP/1.1`	*Client: request line*
`host: www.aol.com`	*Client: required HTTP/1.1 HOST header*
	Client: empty line terminates headers .
`HTTP/1.0 200 OK`	Server: response line
`MIME-Version: 1.0`	*Server: followed by five response headers*
`Date: Mon, 08 Jan 2001 04:59:42 GMT`	
`Server: NaviServer/2.0 AOLserver/2.3.3`	
`Content-Type: text/html`	*Server: expect HTML in the response body*
`Content-Length: 42092`	*Server: expect 42,092 bytes in the resp body*
	Server: empty line ("`\r\n`") terminates hdrs
`<html>`	*Server: first HTML line in response body*
`...`	*Server: 766 lines of HTML not shown.*
`</html>`	*Server: last HTML line in response body*
`Connection closed by foreign host.`	*Server: closes connection*
`unix>`	*Client: closes connection and terminates*

HTTP Requests

- HTTP request is a *request line*, followed by zero or more *request headers*

- Request line:
  - `<method> <uri> <version>`
  - `<version>` is HTTP version of request (`HTTP/1.0` or `HTTP/1.1`)
  - `<uri>` is typically URL for proxies, URL suffix for servers
  - `<method>` is either `GET, POST, OPTIONS, HEAD, PUT, DELETE,` or `TRACE`

HTTP Requests

- HTTP methods:
  - `GET`: Retrieve static or dynamic content
    - Arguments for dynamic content are in URI
    - Workhorse method (99% of requests)
  - `POST`: Retrieve dynamic content
    - Arguments for dynamic content are in the request body
  - `OPTIONS`: Get server or file attributes
  - `HEAD`: Like `GET` but no data in response body
  - `PUT`: Write a file to the server
  - `DELETE`: Delete a file on the server
  - `TRACE`: Echo request in response body
    - Useful for debugging

HTTP Responses

- HTTP response is a *response line* followed by zero or more *response headers*
- Response line:

```
<version> <status code> <status msg>
```

  - <version> is HTTP version of the response
  - <status code> is numeric status
  - <status msg> is corresponding English text
    - 200 OK Request was handled without error
    - 403 Forbidden Server lacks permission to access file
    - 404 Not found Server couldn't find the file
- Response headers: `<header name>: <header data>`
  - Provide additional information about response
  - `Content-Type:` MIME type of content in response body
  - `Content-Length:` Length of content in response body

REST

- Representational State Transfer (REST) is a design pattern.
- A style of software architecture for distributed hypermedia systems such as the World Wide Web
- A collection of network architecture principles which outline how **resources** are defined and addressed.
- It is a certain approach to creating Web Services.
- To understand the REST design pattern, let's look at an example (learn by example).

Example:
Airline Reservation Service

- Suppose that an airline wants to create a telephone reservation system for customers to call in and make flight reservations.
- The airline wants to ensure that its premier members get immediate service, its frequent flyer members get expedited service, and all others get regular service.
- There are two main approaches to implementing the reservation service...

Approach 1
"Press 1 for Premier, Press 2 for..."

The airline provides a single telephone number.

Upon entry into the system a customer encounters an automated message, "Press 1 if you are a premier member, press 2 if you are a frequent flyer, press 3 for all others."

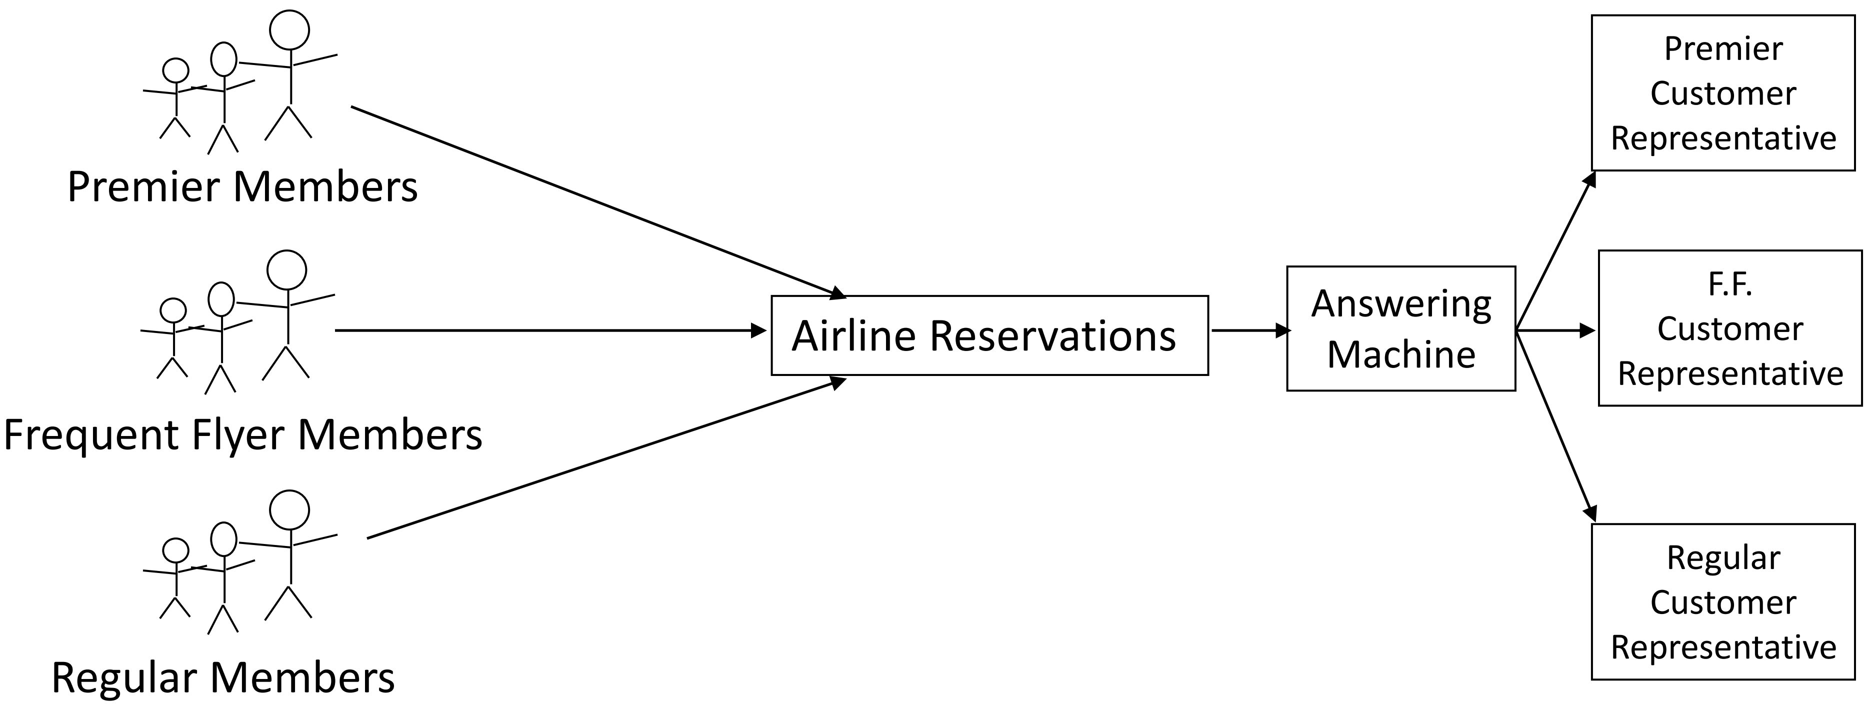

Approach 2
Telephone Numbers are Cheap! Use Them!

The airline provides several telephone numbers - one number for premier members, a different number for frequent flyers, and still another for regular customers.

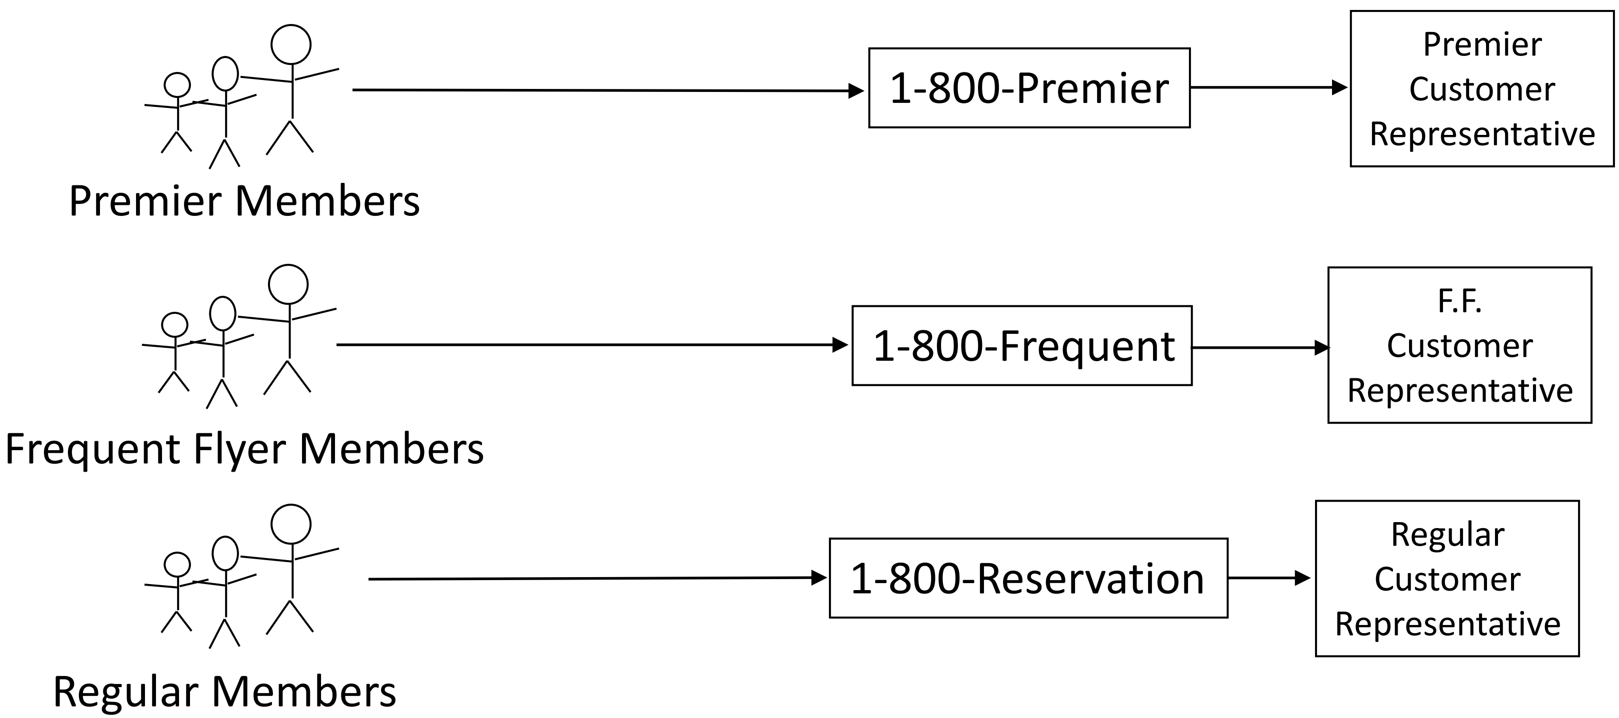

Discussion

- In Approach 1 the answering machine introduces an extra delay, which is particularly annoying to premier members. (Doesn't everyone hate those answering systems)
- With Approach 2 there is no intermediate step. Premier members get instant pickup from a customer service representative. Others may have to wait for an operator.

Web-Based Reservation Service

- Suppose now the airline (kings-air.com) wants to provide a Web reservation service for customers to make flight reservations through the Web.
- Just as with the telephone service, the airline wants to ensure that its premier members get immediate service, its frequent flyer members get expedited service, all others get regular service.
- There are two main approaches to implementing the Web reservation service. The approaches are analogous to the telephone service...

Approach 1
One-Stop Shopping

The airline provides a single URL. The Web service is responsible for examining incoming client requests to determine their priority and process them accordingly.

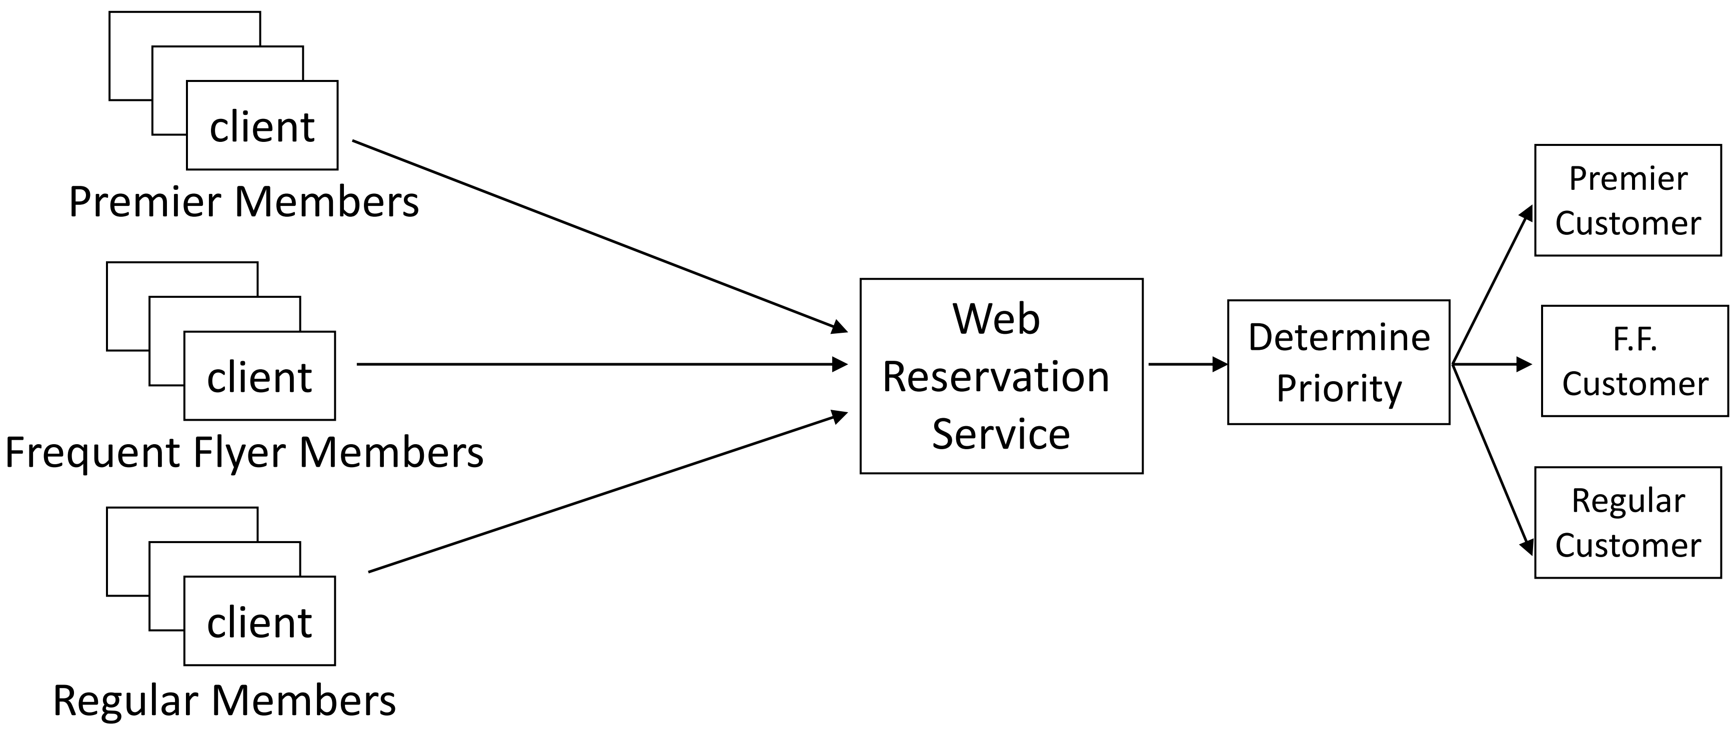

Approach 1 Disadvantages

- There is currently no industry accepted practice (rules) for expressing priorities, so rules would need to be made. The clients must learn the rule, and the Web service application must be written to understand the rule.
- This approach is based upon the incorrect assumption that a URL is "expensive" and that their use must be rationed.
- The Web service is a central point of failure. It is a bottleneck. Load balancing is a challenge.
- It violates Tim Berners-Lee Web Design, Axiom 0 (see next slide).

Web Design, Axiom 0
(Tim Berners-Lee, director of W3C)

- Axiom 0: all resources on the Web must be uniquely identified with a URI.

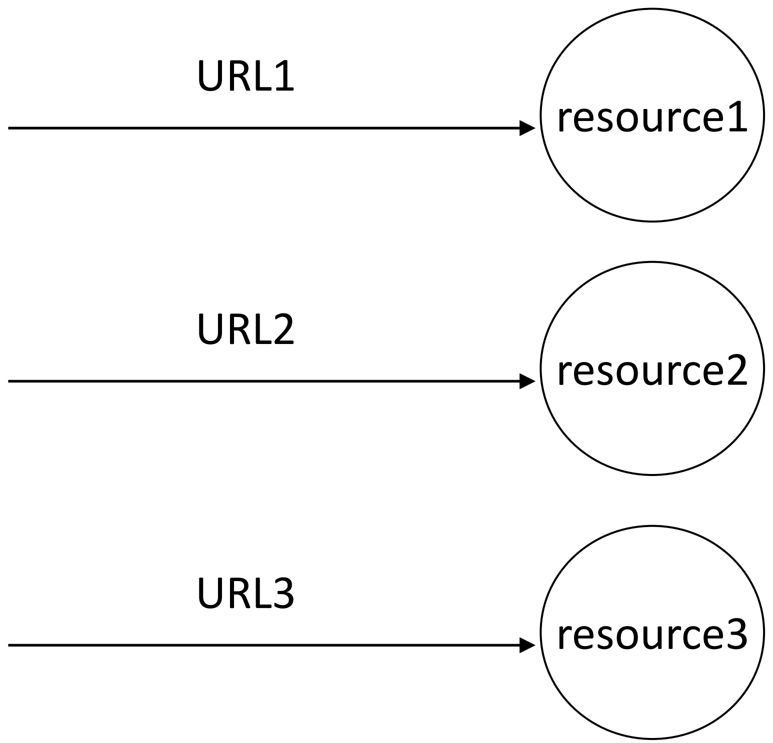

Approach 2:
URLs are Cheap! Use Them!

The airline provides several URLs - one URL for premier members, a different URL for frequent flyers, and still another for regular customers.

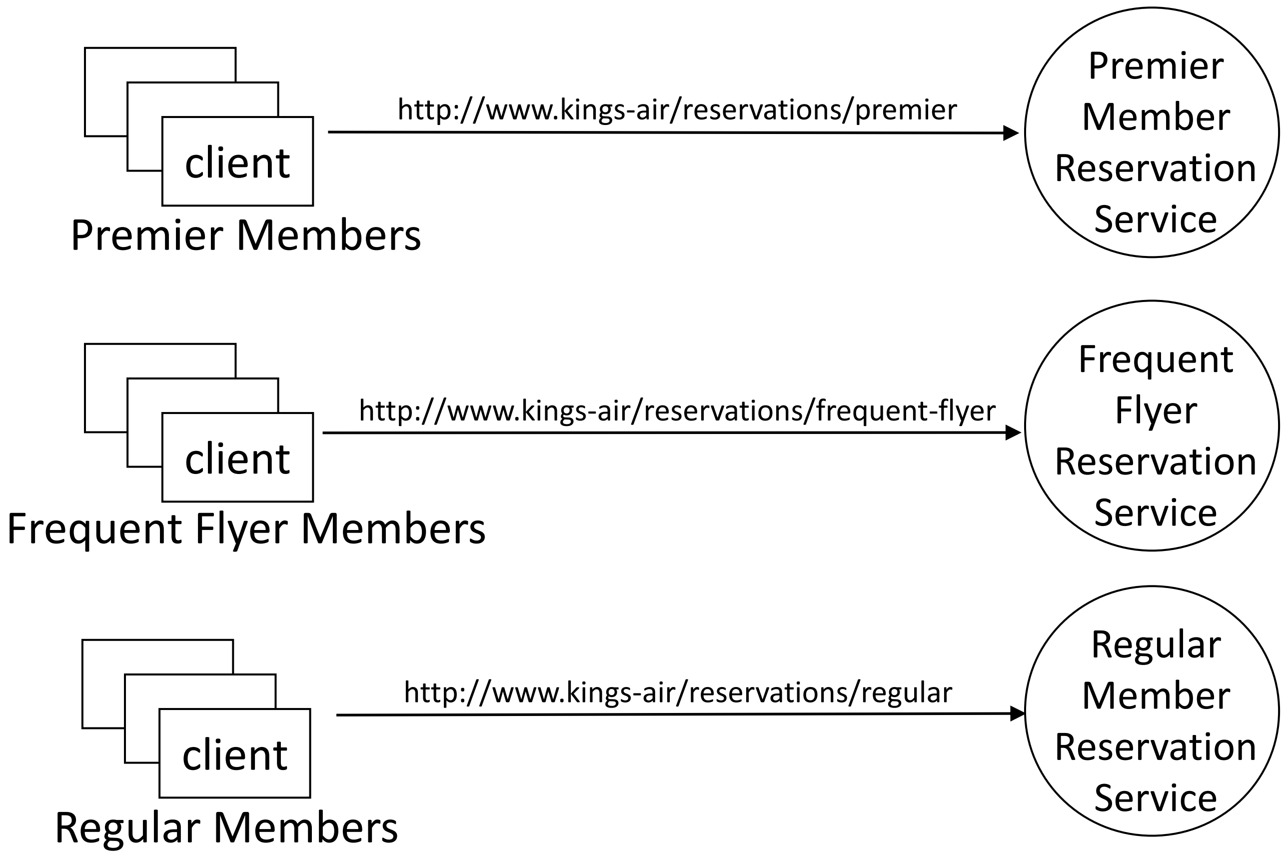

Approach 2 Advantages

- The different URLs are discoverable by search engines and UDDI registries.
- It's easy to understand what each service does simply by examining the URL, *i.e.,* it exploits the Principle of Least Surprise.
- There is no need to introduce rules. Priorities are elevated to the level of a URL. "What you see is what you get."
- It's easy to implement high priority - simply assign a fast machine at the premier member URL.
- There is no bottleneck. There is no central point of failure.
- Consistent with Axiom 0.

Recap

- We have looked at a reservation service.
- We have seen a telephone-based version and a Web-based version of the reservation service.
- With each version we have seen two main approaches to implementing the service.
- Which approach is the REST design pattern and which isn't? See the following slides.

This Ain't the REST Design Pattern

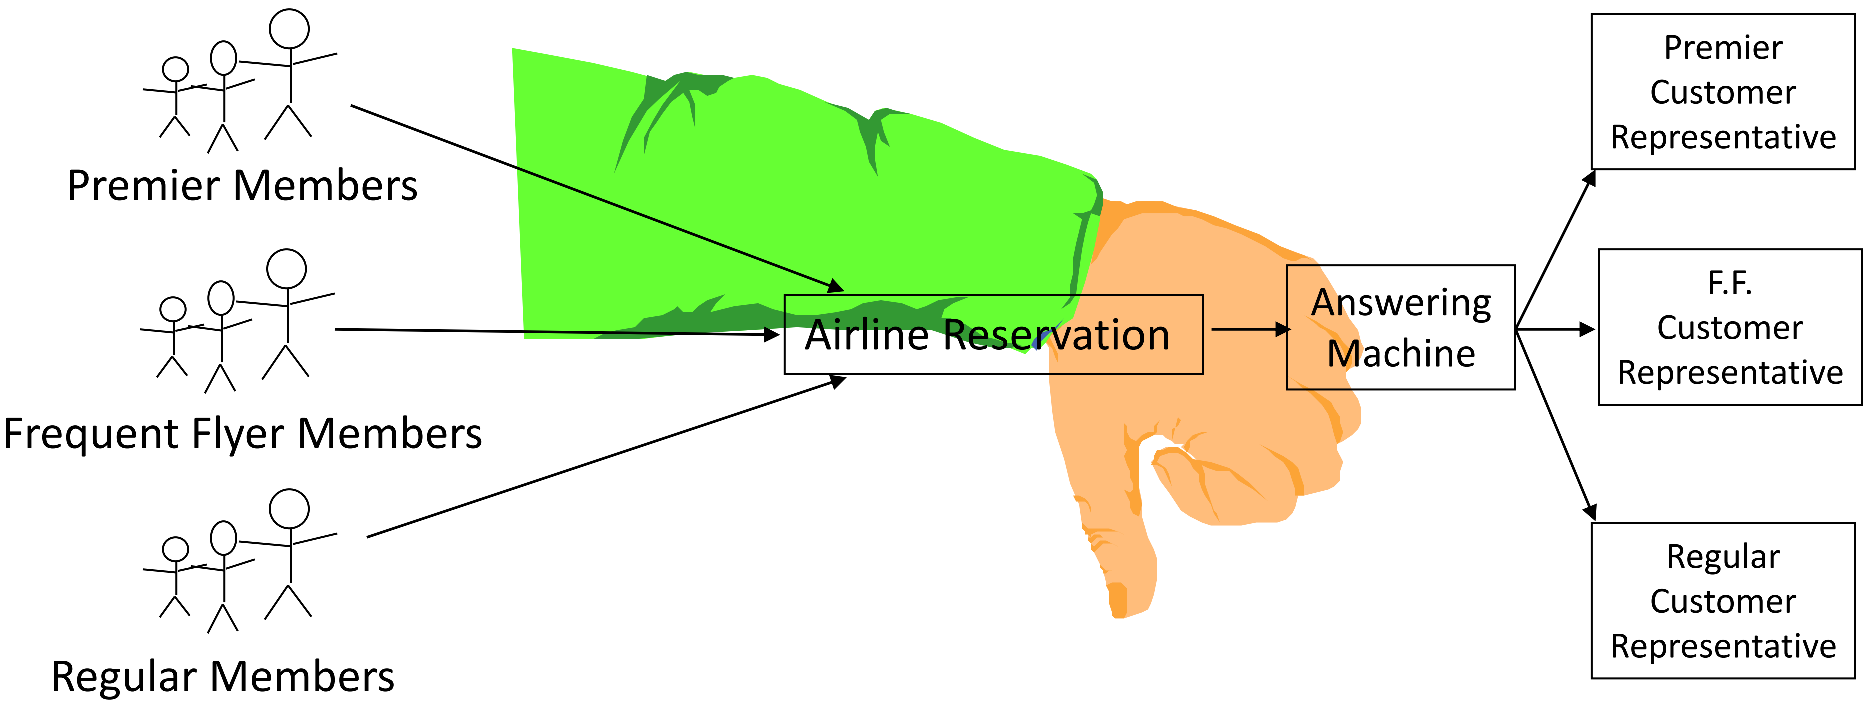

This is the REST Design Pattern

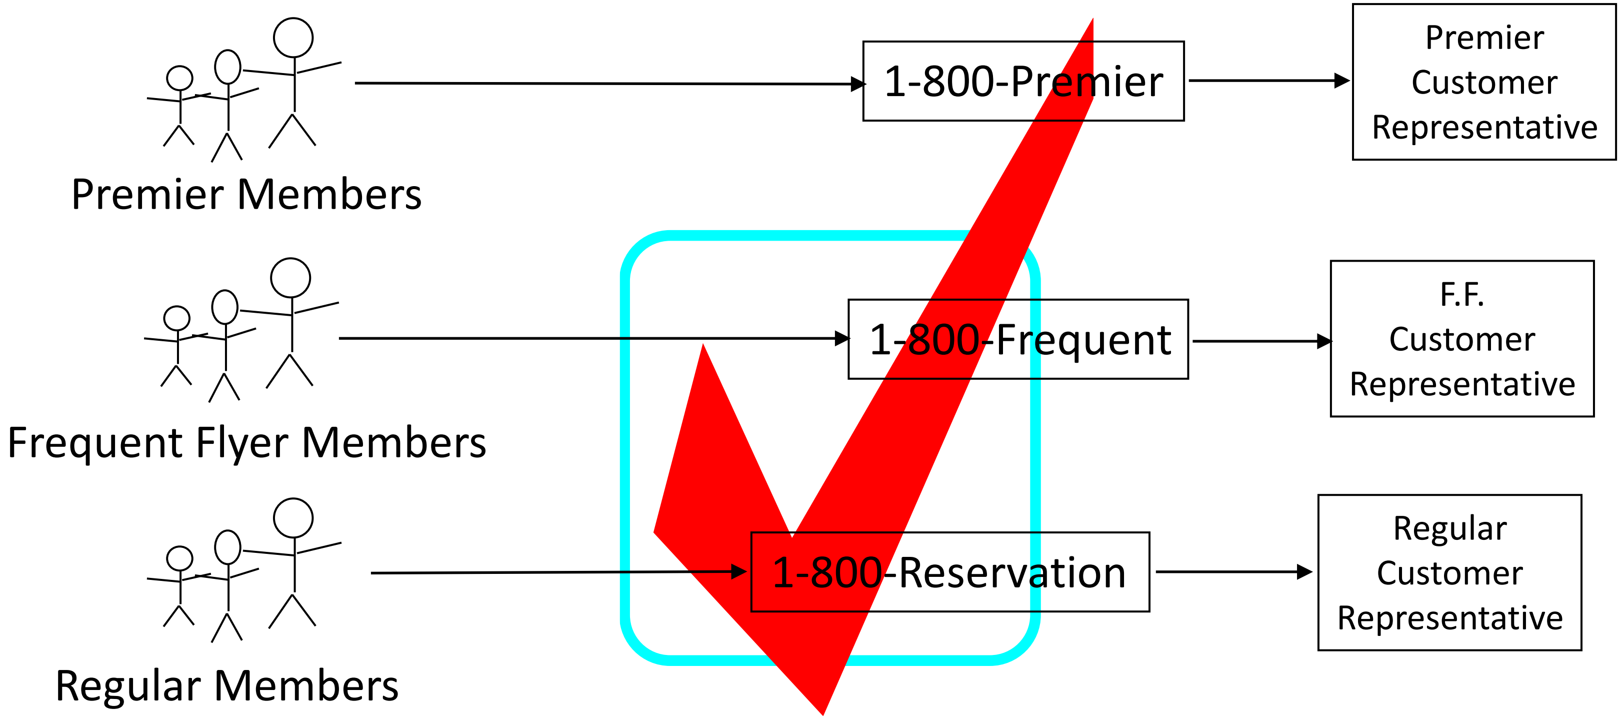

This ain't the
REST Design Pattern

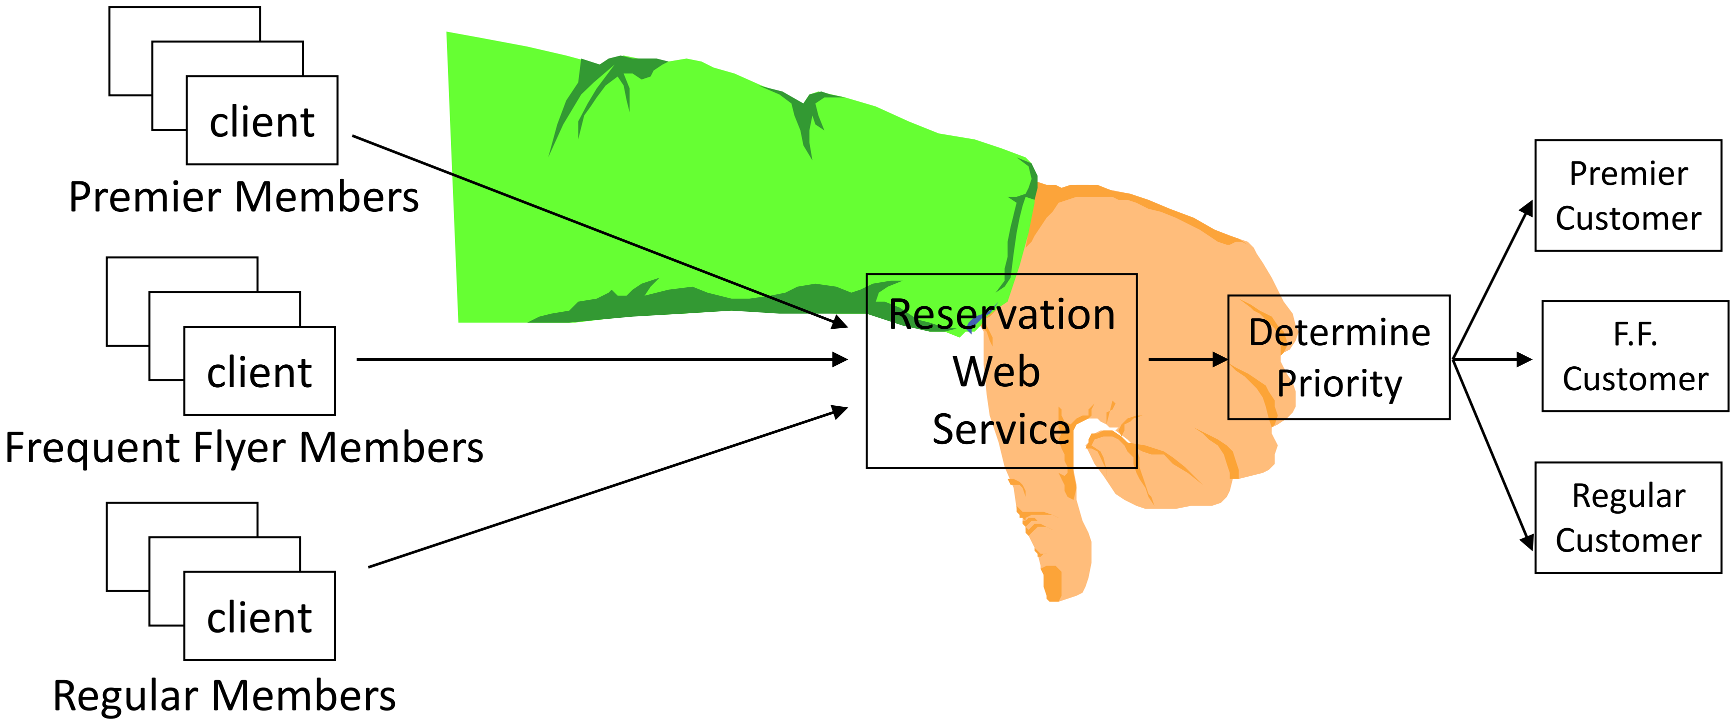

This is the REST Design Pattern

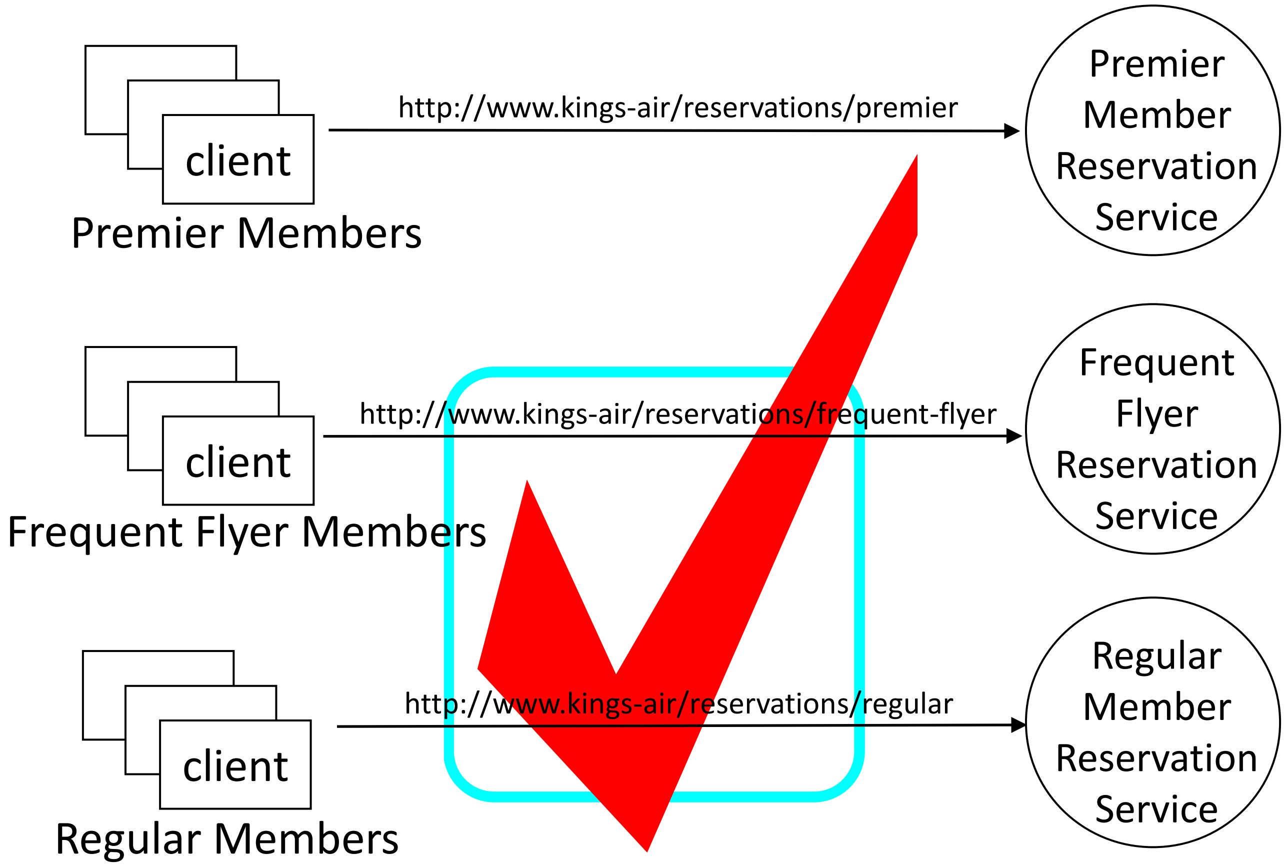

The Three Fundamental Aspects of the REST Design Pattern

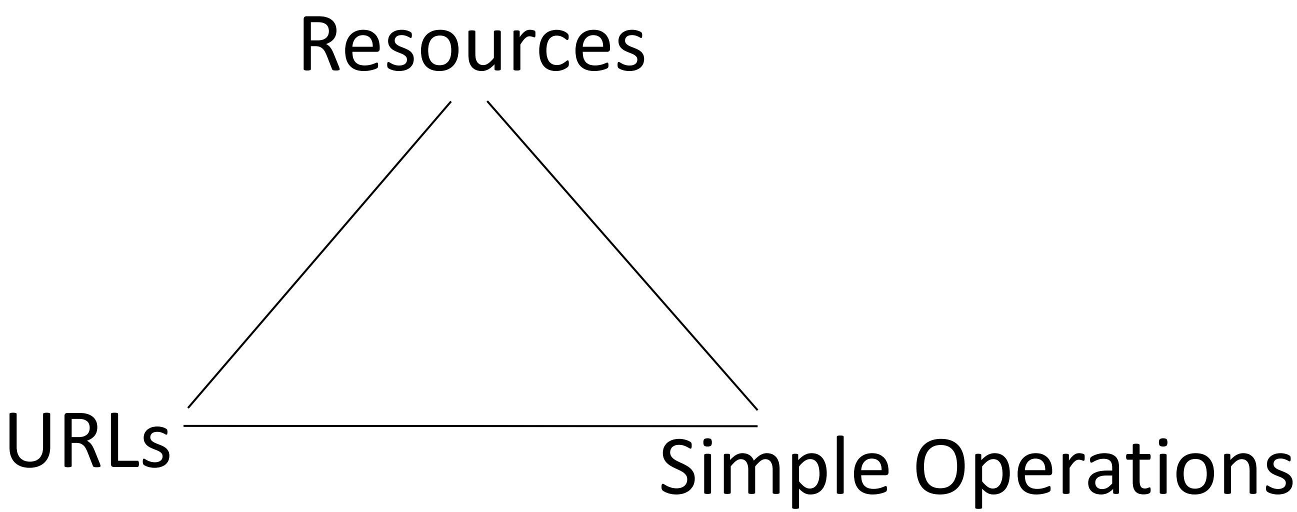

In this tutorial we discussed how Resources and URLs are fundamental to REST. In a follow up tutorial we will discuss how Simple Operations are also fundamental to REST.

REST & HTTP

- The motivation for REST was to capture the characteristics of the Web which made the Web successful
  - URI Addressable resources
  - HTTP Protocol
  - Make a Request – Receive Response – Display Response
- Exploits the use of the HTTP protocol beyond HTTP POST and HTTP GET
  - HTTP PUT, HTTP DELETE

REST

- REST is not a standard
  - is an architectural style
- But it uses several standards:
  - HTTP
  - URL
  - XML/HTML/GIF/JPEG/etc (Resource Representations)
  - text/xml, text/html, image/gif, image/jpeg, etc (Resource Types, MIME Types)

REST Main Concepts

Nouns (Resources)
unconstrained
i.e., http://example.com/employees/12345

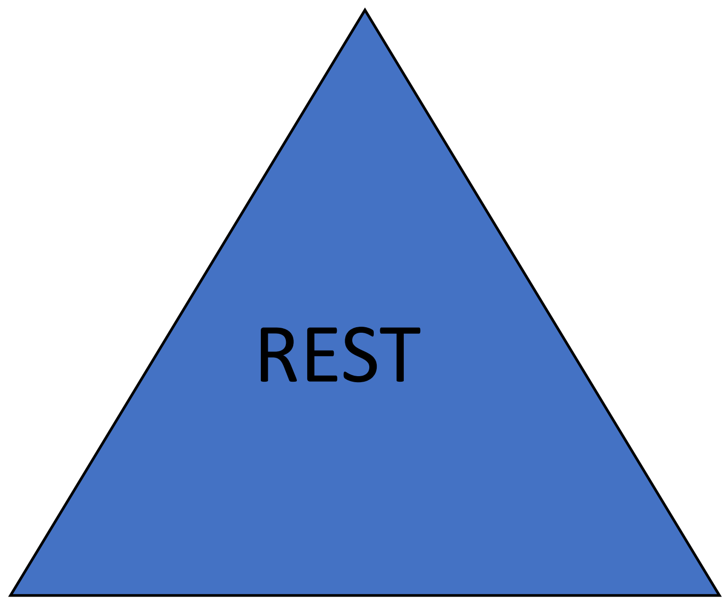

Verbs
constrained
i.e., GET

Representations
constrained
i.e., XML

Resources

- The key abstraction of information in REST is a resource
- A resource is a conceptual mapping to a set of entities
  - Any information that can be named can be a resource: a document or image, a temporal service (e.g., “today's weather in Berlin”), a collection of other resources, a non-virtual object (e.g., a person), etc.
- Represented with a global identifier (URI in HTTP)
  - http://www.boeing.com/aircraft/747

Naming Resources

- REST uses URI to identify resources

  - http://localhost/books/
  - http://localhost/books/ISBN-0011
  - http://localhost/books/ISBN-0011/authors

  - http://localhost/classes
  - http://localhost/classes/cs2650
  - http://localhost/classes/cs2650/students

- As you traverse the path from more generic to more specific, you are navigating the data

Verbs

- Represent the actions to be performed on resources

- HTTP GET
- HTTP POST
- HTTP PUT
- HTTP DELETE

HTTP GET

- How clients ask for the information they seek

- Issuing a GET request transfers the data from the server to the client in some representation

- GET http://localhost/books
  - Retrieve all books

- GET http://localhost/books/ISBN-0011021
  - Retrieve book identified with ISBN-0011021

- GET http://localhost/books/ISBN-0011021/authors
  - Retrieve authors for book identified with ISBN-0011021

HTTP PUT & POST

- HTTP POST creates a resource
- HTTP PUT updates a resource

- POST http://localhost/books/
  - Content: {title, authors[], …}
  - Creates a new book with given properties

- PUT http://localhost/books/isbn-111
  - Content: {isbn, title, authors[], …}
  - Updates book identified by isbn-111 with submitted properties

Representations

- How data is represented or returned to the client for presentation.

- Two main formats:

  - JavaScript Object Notation (JSON)

  - XML

- It is common to have multiple representations of the same data

Representations

- XML

```
<COURSE>
    <ID>CS2650</ID>
    <NAME>Distributed Multimedia Software</NAME>
</COURSE>
```

- JSON

```
{course
    {id: CS2650}
    {name: Distributed Multimedia Software}
}
```

Thank You

Contact me:

gauravsingal789@gmail.com

Gaurav.singal@nsut.ac.in

www.gauravsingal.in

LinkedIn: https://www.linkedin.com/in/gauravsingal789/

Twitter: https://twitter.com/gaurav_singal

CoAP

- Constrained Application Protocol
  - REST-based web transfer protocol
  - manipulates Web resources using the same methods as HTTP: GET, PUT, POST, and DELETE
  - subset of HTTP functionality re-designed for low power embedded devices such as sensors (for IoT and M2M)

CoAP

- TCP overhead is too high and its flow control is not appropriate for short-lived transactions
- UDP has lower overhead and supports multicast

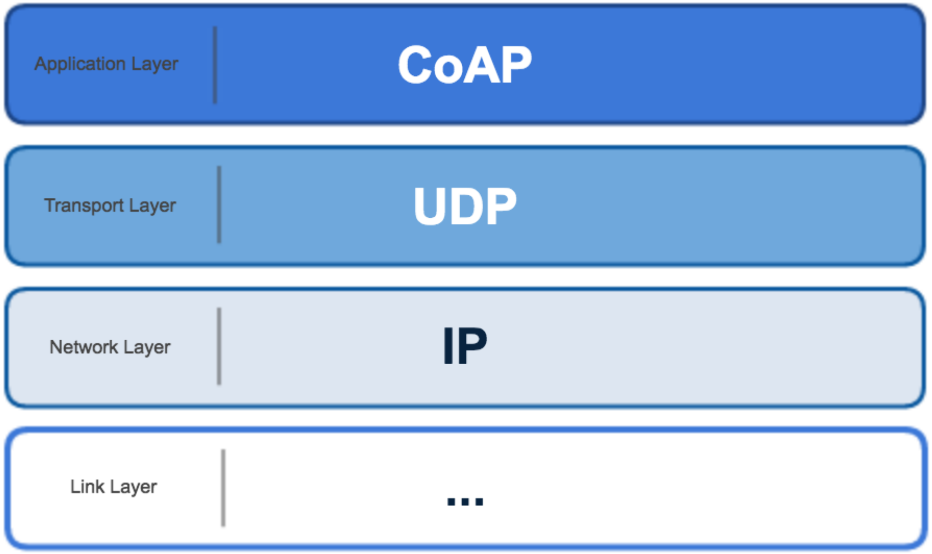

CoAP

- Four message types:
  - **Confirmable** – requires an ACK
  - **Non-confirmable** – no ACK needed
  - **Acknowledgement** – ACKs a Confirmable
  - **Reset** - indicates a Confirmable message has been received but context is missing for processing

CoAP

- CoAP provides reliability without using TCP as transport protocol
- CoAP enables asynchronous communication
  - e.g., when CoAP server receives a request which it cannot handle immediately, it first ACKs the reception of the message and sends back the response in an off-line fashion
- Also supports multicast and congestion control

CoAP

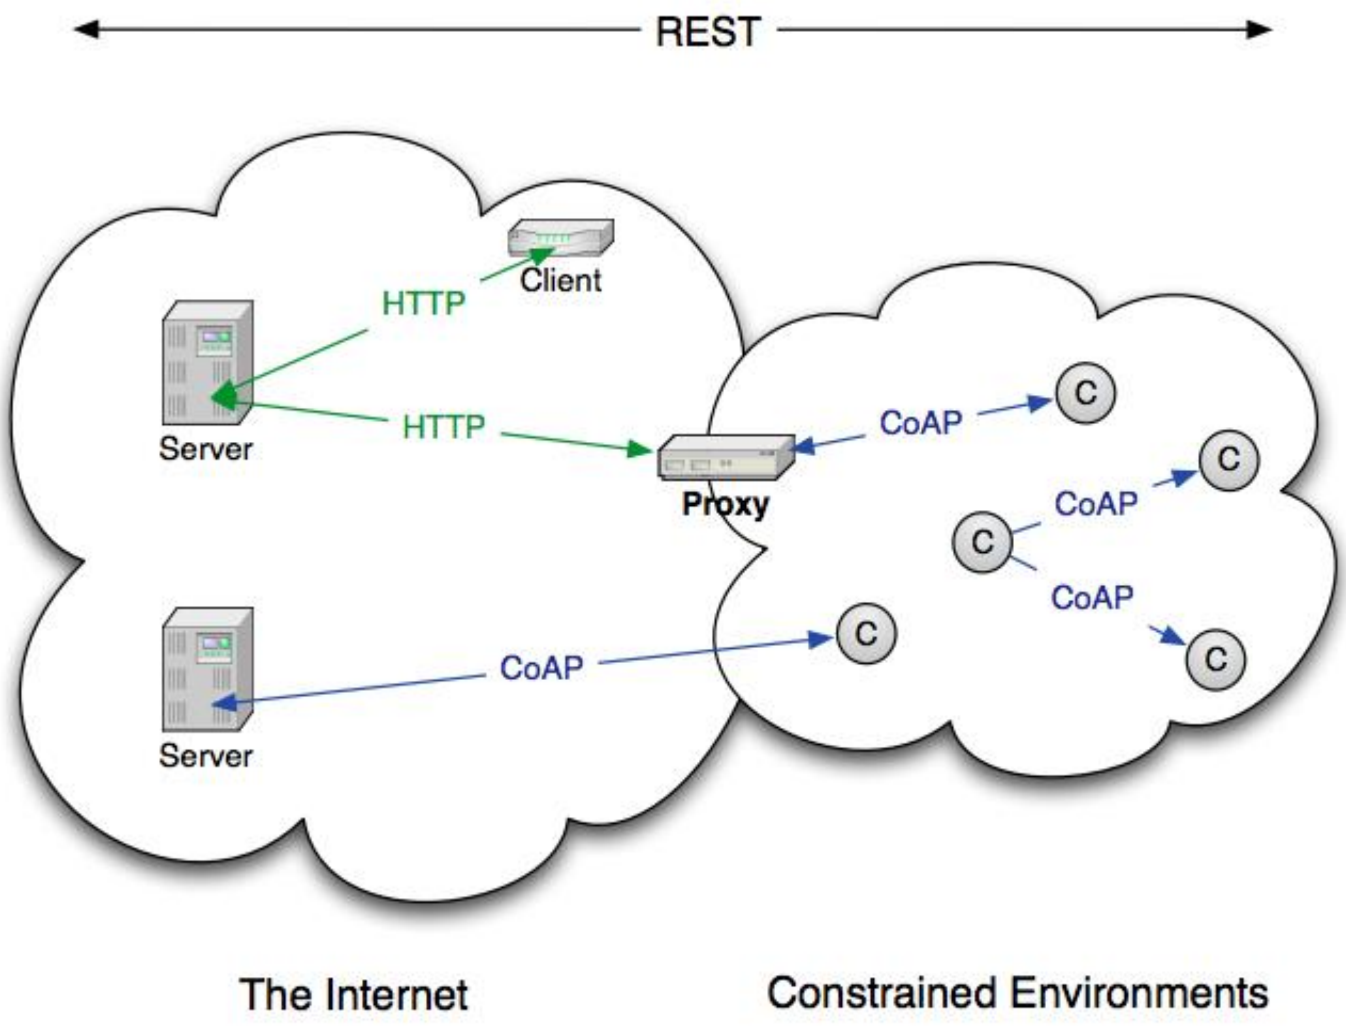
REST
Client
HTTP
Server
HTTP
Proxy
CoAP
C
C
CoAP
C
CoAP
C
Server
CoAP
C
The Internet
Constrained Environments

What CoAP Is

- CoAP is
  - A RESTful protocol
  - Both synchronous and asynchronous
  - For constrained devices and networks
  - Specialized for M2M applications
  - Easy to proxy to/from HTTP

BREAK

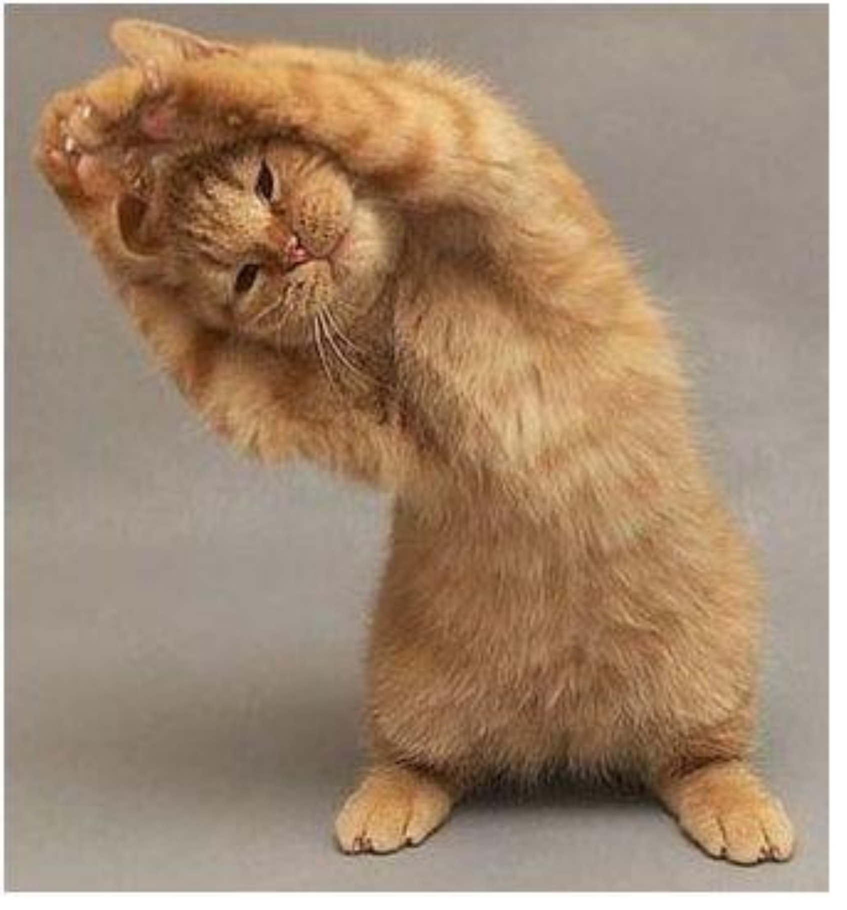

MQTT

- Message Queuing Telemetry Transport
- In a nutshell, MQTT consist of three parts:
  - Broker
  - Subscribers
  - Publishers

MQTT

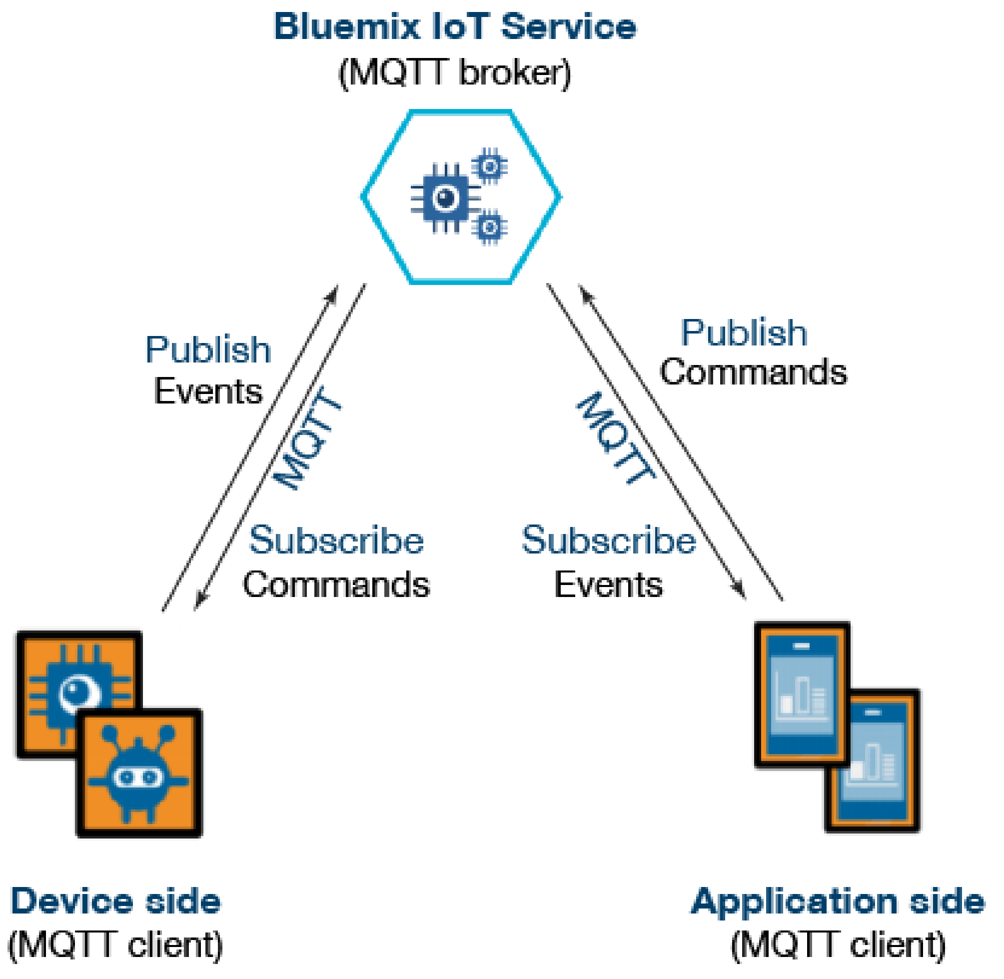

Bluemix IoT Service
(MQTT broker)
Publish
Events
MQTT
Subscribe
Commands
MQTT
Publish
Commands
Subscribe
Events
Device side
(MQTT client)
Application side
(MQTT client)

MQTT

- MQTT was invented by Andy Stanford-Clark (IBM) and Arlen Nipper (Arcom, now Cirrus Link) back in 1999, where their use case was to create a protocol for minimal battery loss and minimal bandwidth connecting oil pipelines over satellite connections. They specified the following goals, which the future protocol should have:
  - Simple to implement
  - Provide a Quality of Service Data Delivery
  - Lightweight and Bandwidth Efficient
  - Data Agnostic
  - Continuous Session Awareness

MQTT

- Built for proprietary embedded systems; now shifting to IoT
- You can send anything as a message; up to 256 MB
- Built for unreliable networks
- Enterprise scale implementations down to hobby projects
- Decouples readers and writers
- Message have a topic, quality of service, and retain status associated with them

Publish/Subscribe Concept

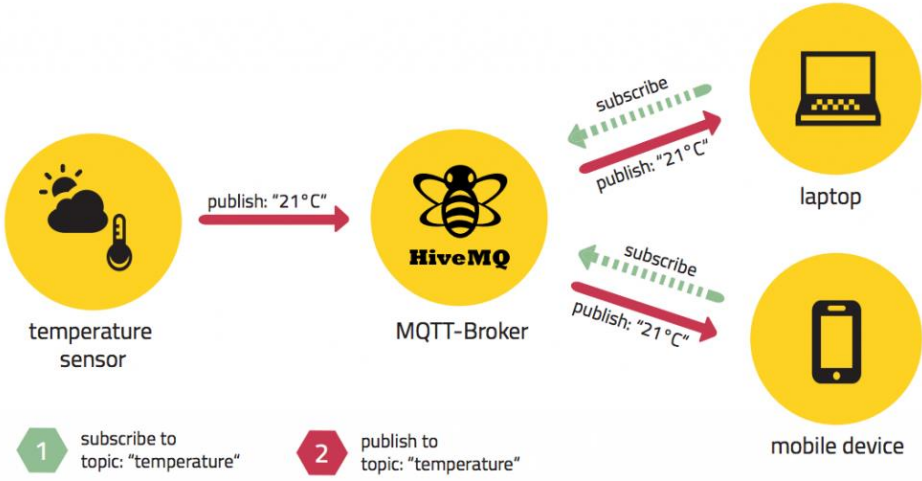

Decoupled in space and time:
The clients do not need each others IP address and port (space) and they do not need to be running at the same time (time).

The broker's IP and port must be known by clients

Namespace hierarchy used for topic filtering

It may be the case that a published message is never consumed by any subscriber

MQTT: Example

- Clients connect to a “Broker”
- Clients subscribe to topics e.g.,
  - client.subscribe(‘toggleLight/1’)
  - client.subscribe(‘toggleLight/2’)
  - client.subscribe(‘toggleLight/3’)
- Clients can publish messages to topics:
  - client.publish(‘toggleLight/1’, ‘toggle’);
  - client.publish(‘toggleLight/2’, ‘toggle’);
- All clients receive all messages published to topics they subscribe to
- **Messages can be anything**
  - Text
  - Images
  - etc.

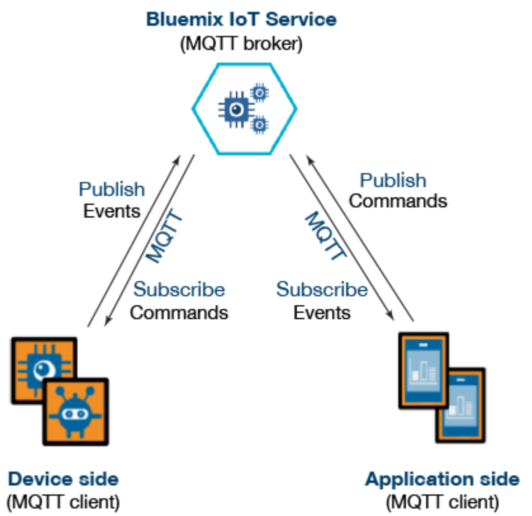

Node.js Example

```
var mqtt = require('mqtt');
var client = mqtt.createClient('<<PortNumber>>', 'm11.cloudmqtt.com', {
        username: '<<UserName>>',
        password: '<<Password>>'
});
client.on('connect', function () { // When connected
        // subscribe to a topic
        client.subscribe('TEMPERATURE_READING', function () {
                // when a message arrives, do something with it
                client.on('message', function (topic, message, packet) {
                        console.log("Received '" + message + "' on '" + topic + "'");
                });
        });
        // publish a message to a topic
        client.publish('SET_TEMPERATURE', '24', function () {
                console.log("Message is published");
        });
});
```

Topics

- Each published data specifies a topic
- Each subscriber subscribed to that topic will receive it
- Topic format:

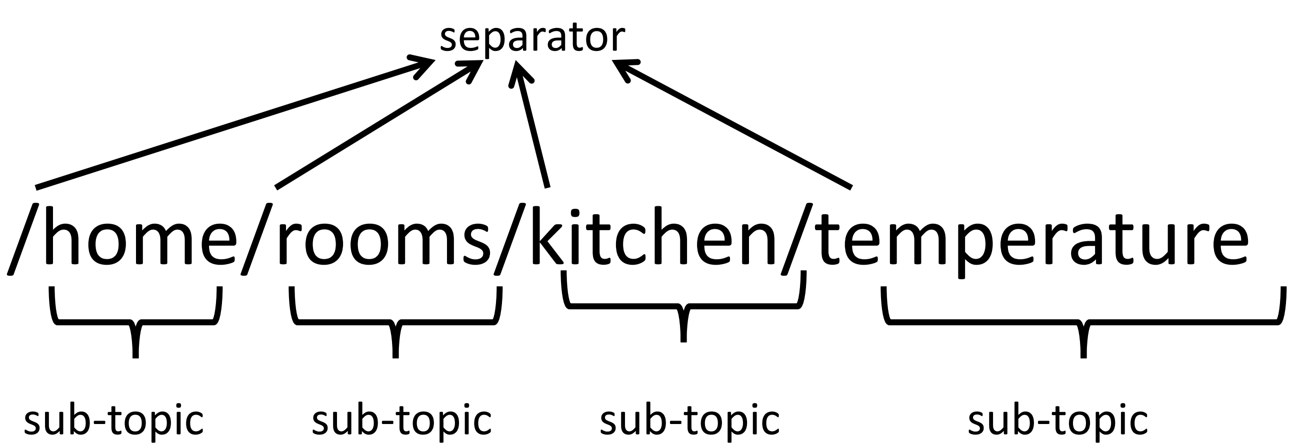

Durable/Transient Subscriptions

- Subscriptions
  - Durable
    - If the subscriber disconnect messages are buffered at the broker and delivered upon reconnection
  - Non-durable
    - Connection lifetime gives subscription lifetime

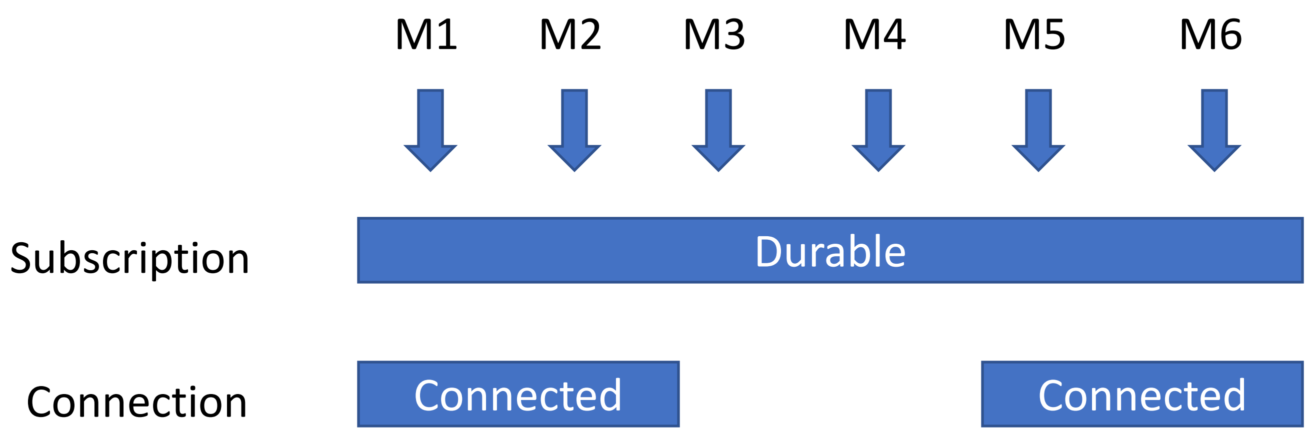

State Retention

- Publications
  - Retained (“persistent” message)
    - The subscriber upon first connection receives the last good publication (i.e., does not have to wait for new publication)
  - One flag set both in the publish packet to the broker and in the published packet to the subscribers
    - Only the most recent persistent message is stored and distributed

Session Aware

- Last Will and Testament (LWT) – topic published upon disconnecting a connection
- Any client can register a LWT
- Anybody subscribing to the LWT topic will know when a certain device (that registered a LWT) disconnected

Protocol Stack

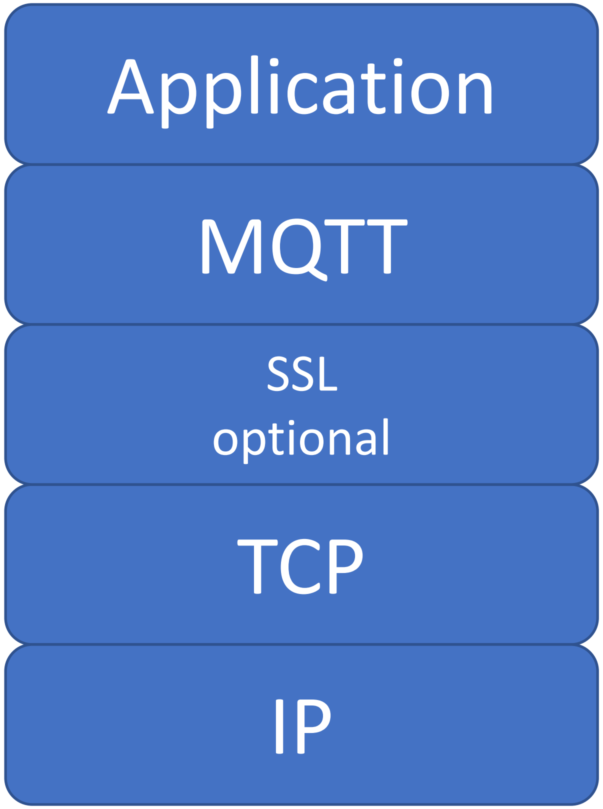

TCP/IP Port: 1883
When running over SSL, TCP/IP port 8883
SSL: Secure Socket Layer (encryption)

Publishing "QoS" (Reliability)

- 0 – unreliable (aka "at most once")
  - OK for continuous streams, least overhead (1 message)
  - "Fire and forget"
  - TCP will still provide reliability

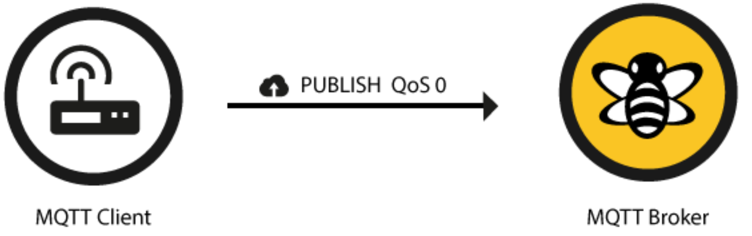

Publishing "QoS" (Reliability)

- 1 – delivery "at least once" (duplicates possible)
  - Used for alarms – more overhead (2 messages)
  - Contains message ID (to match with ACKed message)

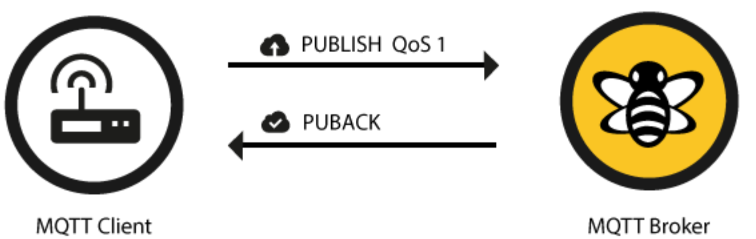

Publishing "QoS" (Reliability)

- 2 – delivery "exactly once"
  - Utmost reliability is important – most overhead (4 messages) and slowest

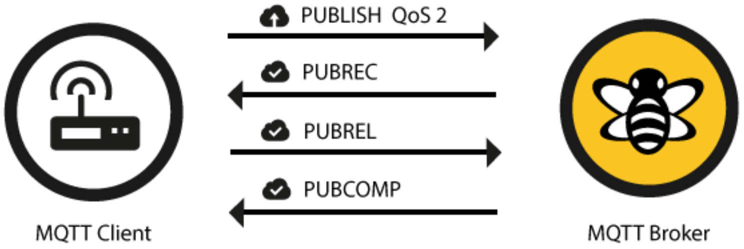

Publishing “QoS” (Reliability)

- Reliability maintained even if the TCP connection breaks (intermittent connections)
- Separate QoS for publishing and for subscribing

MQTT Message Format

Shortest Message is Two Bytes

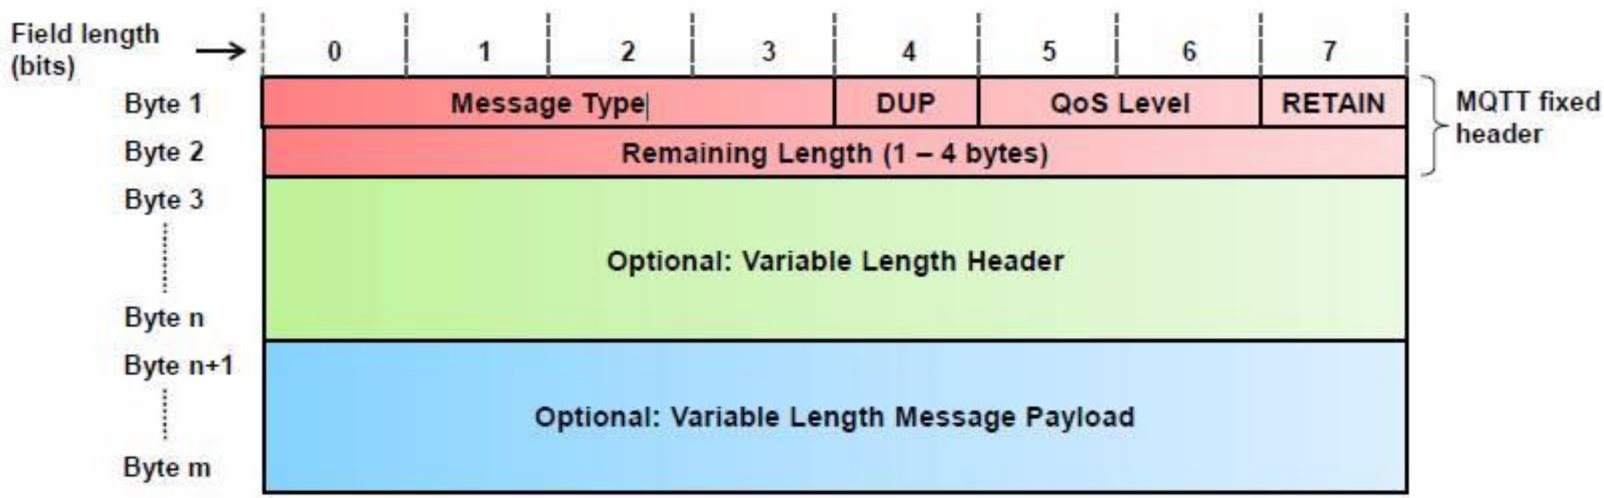

Message Types

Name	Value	Direction of flow	Description
Reserved	0	Forbidden	Reserved
CONNECT	1	Client to Server	Client request to connect to Server
CONNACK	2	Server to Client	Connect acknowledgment
PUBLISH	3	Client to Server or Server to Client	Publish message
PUBACK	4	Client to Server or Server to Client	Publish acknowledgment
PUBREC	5	Client to Server or Server to Client	Publish received (assured delivery part 1)
PUBREL	6	Client to Server or Server to Client	Publish release (assured delivery part 2)
PUBCOMP	7	Client to Server or Server to Client	Publish complete (assured delivery part 3)
SUBSCRIBE	8	Client to Server	Client subscribe request
SUBACK	9	Server to Client	Subscribe acknowledgment
UNSUBSCRIBE	10	Client to Server	Unsubscribe request
UNSUBACK	11	Server to Client	Unsubscribe acknowledgment
PINGREQ	12	Client to Server	PING request
PINGRESP	13	Server to Client	PING response
DISCONNECT	14	Client to Server	Client is disconnecting
Reserved	15	Forbidden	Reserved

Message Types

Message fixed header field	Description / Values	
Message Type	**0: Reserved**	**8: SUBSCRIBE**
	1: CONNECT	**9: SUBACK**
	2: CONNACK	**10: UNSUBSCRIBE**
	3: PUBLISH	**11: UNSUBACK**
	4: PUBACK	**12: PINGREQ**
	5: PUBREC	**13: PINGRESP**
	6: PUBREL	**14: DISCONNECT**
	7: PUBCOMP	**15: Reserved**
DUP	**Duplicate message flag. Indicates to the receiver that this message may have already been received. 1: Client or server (broker) re-delivers a PUBLISH, PUBREL, SUBSCRIBE or UNSUBSCRIBE message (duplicate message).**	
QoS Level	**Indicates the level of delivery assurance of a PUBLISH message. 0: At-most-once delivery, no guarantees, «Fire and Forget». 1: At-least-once delivery, acknowledged delivery. 2: Exactly-once delivery. Further details see MQTT QoS.**	
RETAIN	**1: Instructs the server to retain the last received PUBLISH message and deliver it as a first message to new subscriptions. Further details see RETAIN (keep last message).**	
Remaining Length	**Indicates the number of remaining bytes in the message, i.e. the length of the (optional) variable length header and (optional) payload. Further details see Remaining length (RL).**	

Comparison CoAP & MQTT

Both used in IoT

- CoAP:
  - one-to-one communication
  - UDP/IP
  - unreliable
  - lightweight and easy to implement
- MQTT:
  - many-to-many communication
  - TCP/IP
  - focus on message delivery; reliable
  - higher overheads (protocol data, processing costs)

BREAK